

# Modern Architectures, Profiling and Optimization

Luca Tornatore - I.N.A.F.



Scientific and high performance computing school 2019

Università degli Studi di Trento

# Introduction

High Performance Computing requires, by the name itself, to squeeze as much effectiveness from your code on the machine you run it.

“Optimizing” is, obviously, a key step in this process.

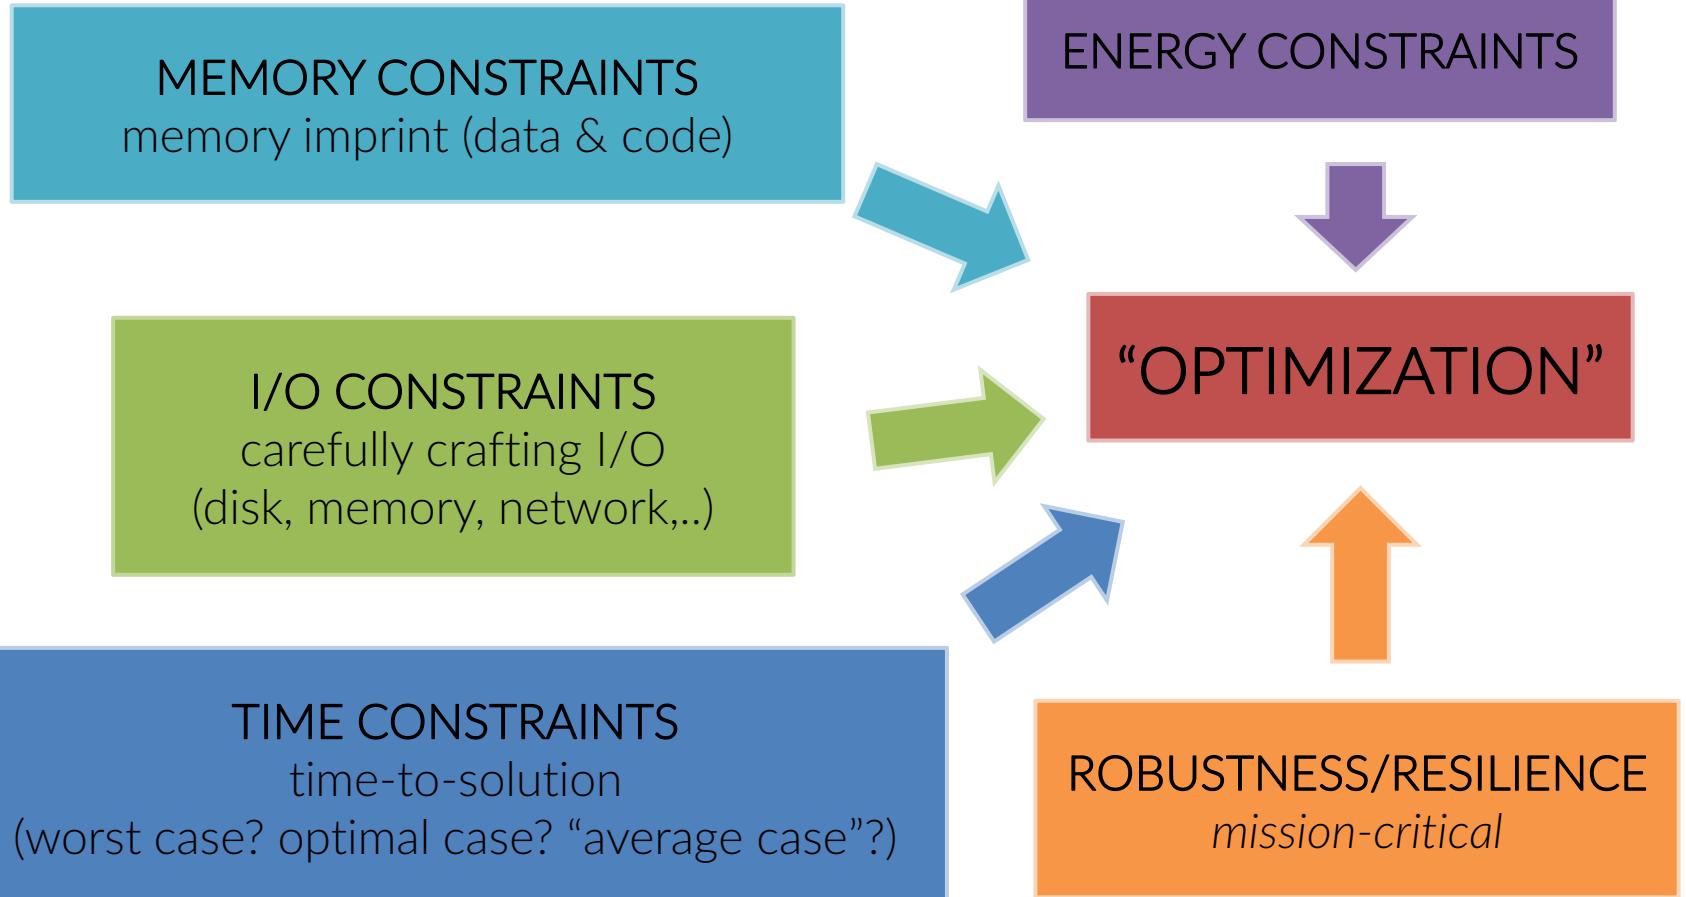
However, both “*optimizing*” and “*effectiveness*” have to be defined. And that is what this lecture is about.

Let’s start with some general comments.

There may be no such thing as an “optimal” code *tout court*

---

Just try to define “optimal”..



*Premature optimization is the root of all evil*  
D. Knuth

You'd better start thinking in terms of “*improved*” code.

- You don't want to hurt your code that honestly does its job: a faster code that gives **incorrect results** is not optimal in *any* way
- You most probably will have to **re-use** that code after some time. And you will need it to be **readable, maintainable** and **well-designed**.
- Others most probably will have to read, understand and re-use that code: a **clean, understandable, non-obscure, non-redundant** code may not be “improved” in some way, but actually improves the quality of your life.

# First steps to consider



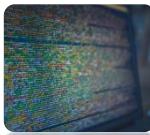
To have a **cleaner** code



To **re-design** its fundamental architecture



To **improve** the workflow, the memory imprint, the resource usage, the time-to-solution, ...



# Dryness

**DRY:**  
Don't  
Repeat  
Yourself

Do neither add **unnecessary** code nor **duplicate** code

**Unnecessary code** increases the amount of needed work

- to maintain the code, either debugging or updating it
- to extend its functionalities

**Duplicated code** increases your *bad technical debt*, that already has a large enough number of sources:

- copy-and-paste programming style
- too much *agile* approach
- hard coding, quick-and-dirty fixes, cargo-cult, sloppiness



# Readability

Readability is a pillar of maintainability.

Maintainability means that software can be extended, upgraded, debugged, fixed.

- **Baseline:** write comments
- **Advanced:** WRITE COMMENTS ! (possibly, use doxygen-like tools)
- **X-advanced:** avoid stupid, obvious, useless, confusing comments

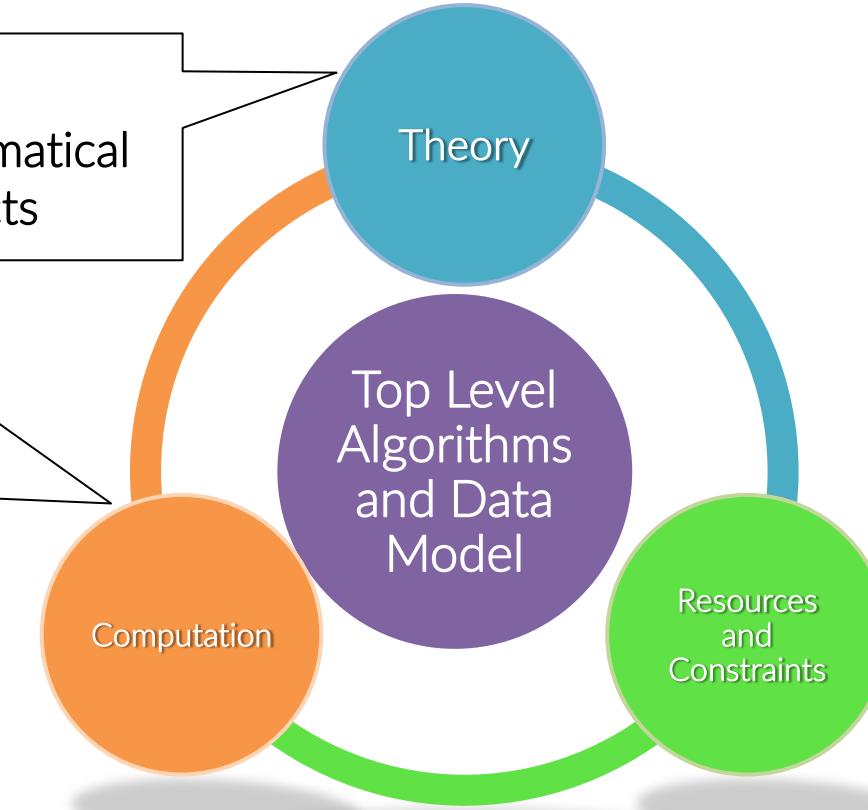
Keep in mind: “*the code is the ultimate man page*”



# Design

- Consult literature
- Understand mathematical and scientific aspects

You're not at the blackboard, but on a digital discrete system. Equations have to *translated*.



- Concurrency & parallelism
- Memory constrs.
- I/O constrs.
- Energy constrs.



Design

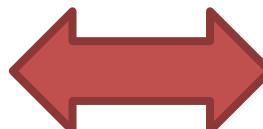
# Design



Highest abstract picture:

- Interacting multi-components
- Top-level DATA MODEL

Sub-systems and modules  
Communication strategies



Modules as producers of complex tasks

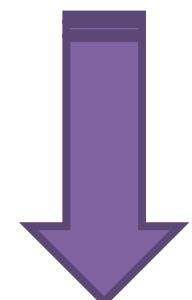
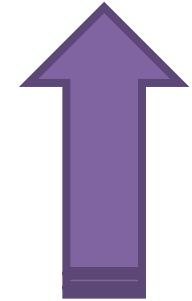
Implementation details

- many algorithms
- concurrency

Logical structure of modules

interactions

Very low-level coding and details



ABSTRACTION

CODING DETAIL



**TESTING** IS part of the design

- Unit test  
*separately stress each unit of the code*
- Integration test  
*stress the integrated behavior of all the units together*
- System test



# Design

**VALIDATION** and **VERIFICATION** ARE part of the design

- Validation ensures that the code does what it was meant to do  
*(all modules are designed accordingly to design specifications)*
- Verification ensures that the codes does what it does correctly  
*(black-box testing against test-cases, ...)*



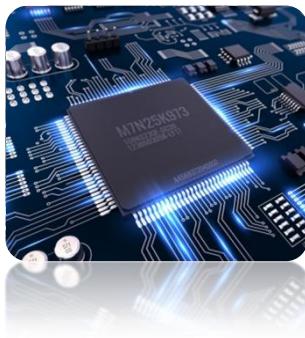
Improve

# Improve the code

**Improving** the workflow, the memory imprint, the resource usage, the time-to-solution, ...

Well, that is the focus of this day :)

# Outline



Modern  
Architecture &  
Optimization



Profiling



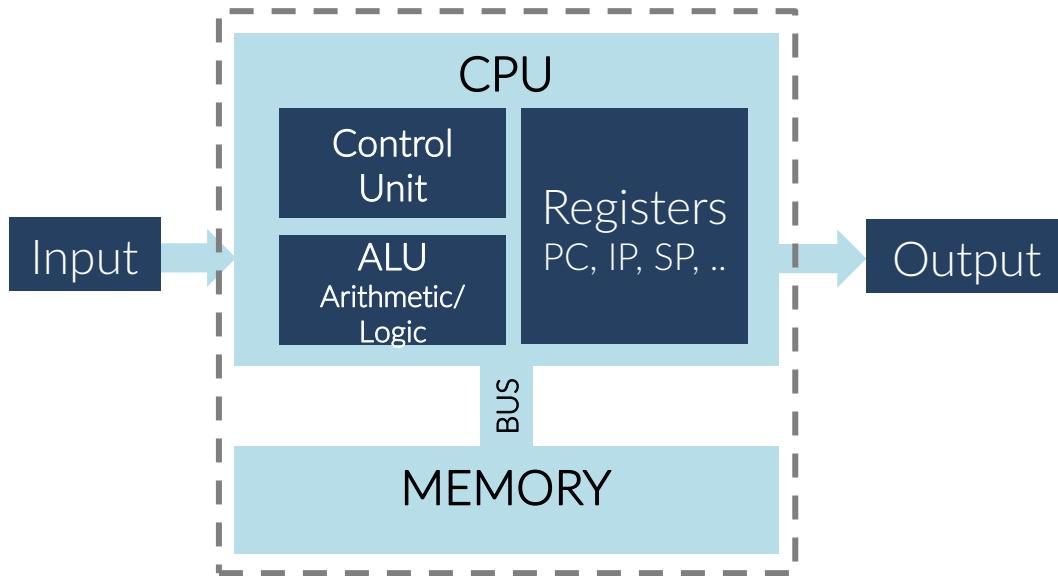
Debugging



# Once upon a time..

.. computer architecture was much simpler than today.

The exact reasons why it was so – or, better, why nowadays it is much more complex – will be discussed tomorrow.

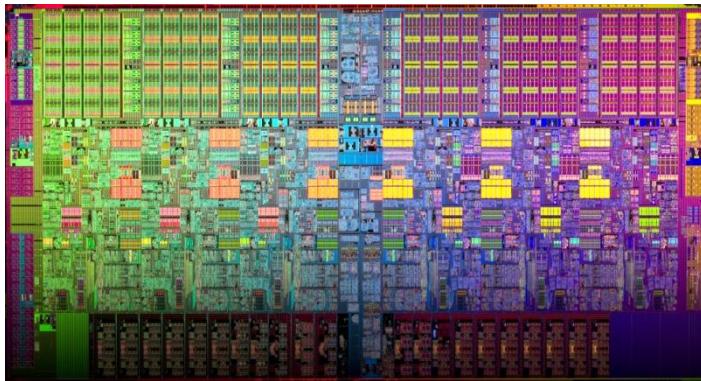


In the **Von Neumann architecture** (still taught as the fundamental model)

- there is only 1 processing unit
- 1 instruction is executed at a time
- memory is “flat”:
  - access on any location has always the same cost
  - access to memory has the same cost than op execution



# While today...



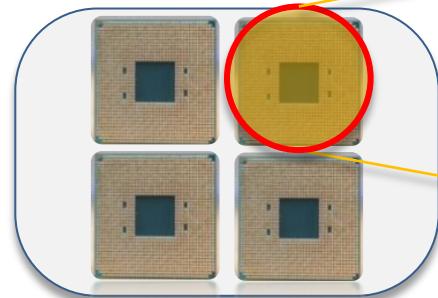
6-cores Intel Xeon e5600

- there *many* processing units
- many instructions can be executed at a time
- many data can be processed at a time
- “instructions” are internally broken down in many simpler  $\mu$ ops that are pipelined
  - different instructions could have quite different cost
- memory is strongly not “flat”: (\*)
  - there is a strong memory hierarchy
  - access memory can have *very* different costs depending on the location
  - accessing RAM is way more costly than performing a  $\mu$ op or reading an internal register

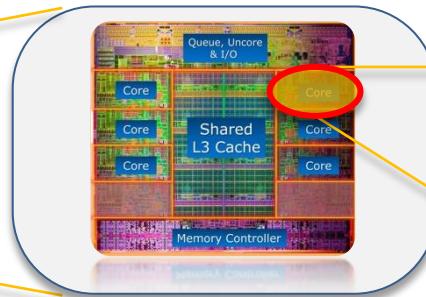
*In the next slides we'll go through all this features in chronological order, as they developed in time*



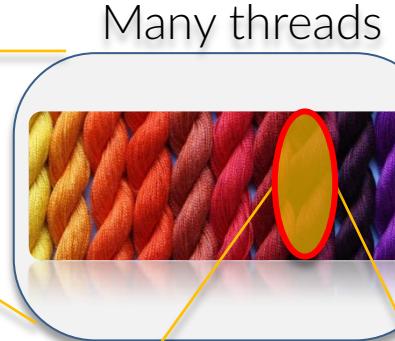
# Overview



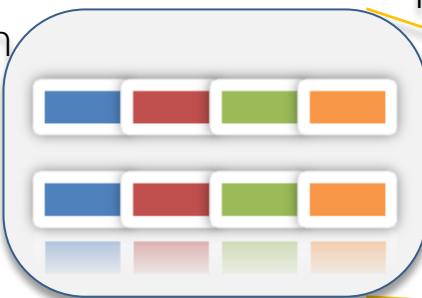
sockets



Many cores  
per socket



Many threads



Vectorisation



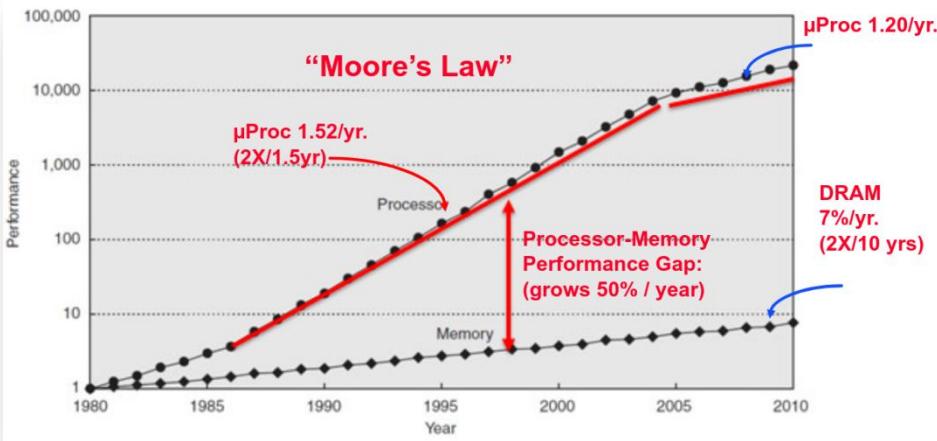
Pipelining



Superscalarity



# Early 90s: CPU get faster than memory



Processor	Alpha 21164	
Machine	AlphaServer 8200	
Clock Rate	300 MHz	
Memory Performance	Latency	Bandwidth
I Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
D Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
L2 Cache (96KB on chip)	20 ns (6 clocks)	4800 MB/sec
L3 Cache (4MB off chip)	26 ns (8 clocks)	960 MB/sec
Main Memory Subsystem	253 ns (76 clocks)	1200 MB/sec
Single DRAM component	$\approx 60\text{ns}$ (18 clocks)	$\approx 30\text{--}100$ MB/sec

Taken from a 1997 paper

The CPU may spend more time waiting for data coming from RAM than executing ops. That is part of the so called “memory wall”. What is the solution ?



# CPU-RAM gap

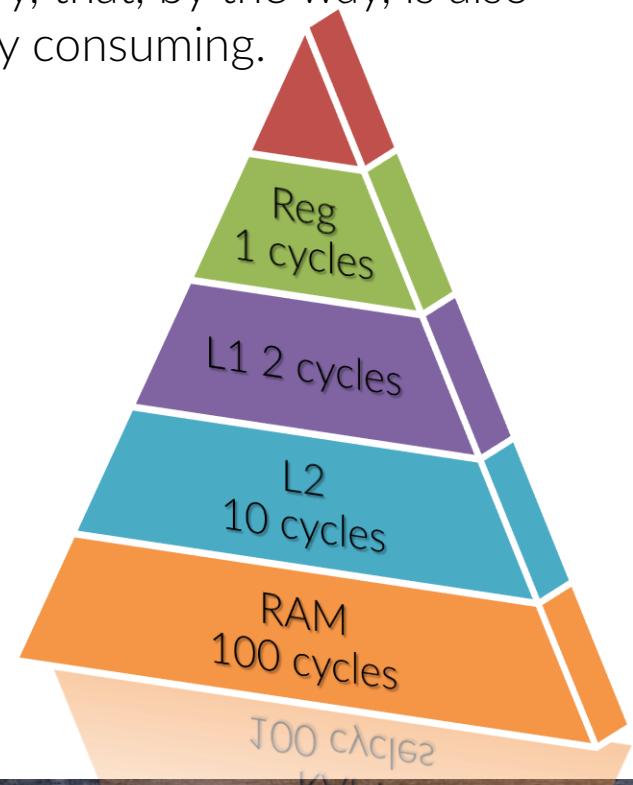
The solution is to equip your CPU with a *faster* memory, that, by the way, is also way more costly (in terms of \$/B) and way more energy consuming.

Furthermore, to be faster it ought to be *extremely* closer.

All in all, the new memory that will be called *cache*, will be much *smaller* than RAM.

The cache itself has a hierarchy:

- *Level-I* is inside each core.
- *Level-II* is also inside the core, or may be shared by more cores.
- *Level-III* is inside the CPU, shared by many cores.





# CPU-RAM gap

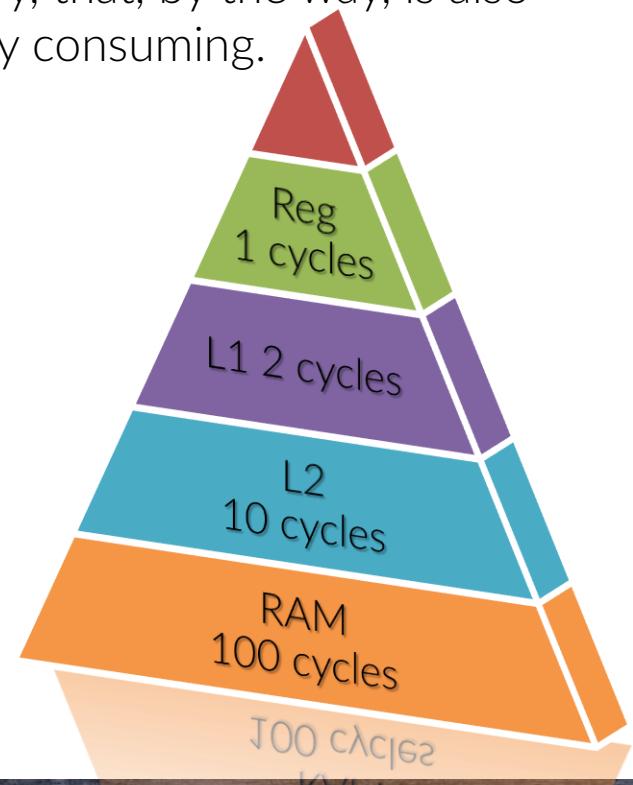
The solution is to equip your CPU with a *faster* memory, that, by the way, is also way more costly (in terms of \$/B) and way more energy consuming.

Furthermore, to be faster it ought to be *extremely* closer.

All in all, the new memory that will be called *cache*, will be much *smaller* than RAM.

The cache itself has a hierarchy:

- *Level-I* is inside each core.
- *Level-II* is also inside the core, or may be shared by more cores.
- *Level-III* is inside the CPU, shared by many cores.



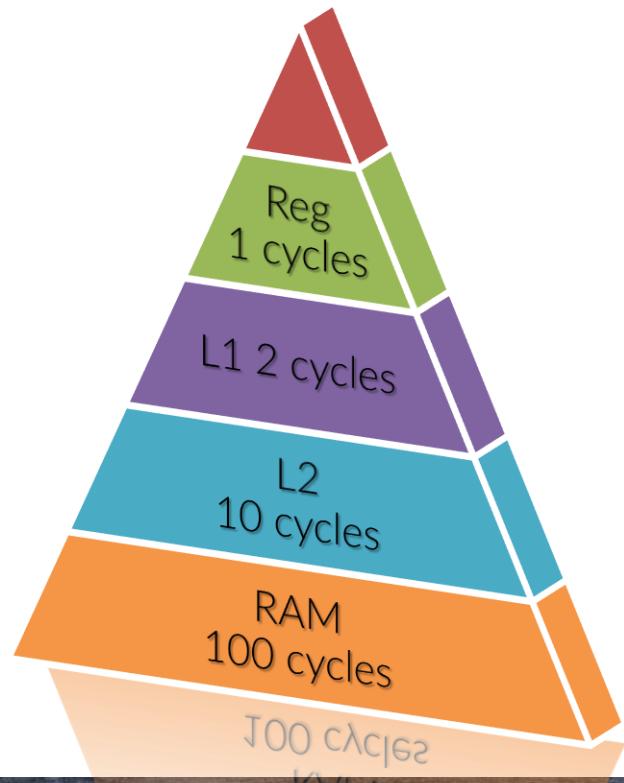


# CPU-RAM gap

L1 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times \text{RAM}_{\text{cost}}$$

- 100% L1 hit → 2 cycles
  - 99% L1 hit → 3 cycles
  - 97% L1 hit → 5 cycles
- } 50% to 150% slower





# CPU-RAM gap

L1 cache + RAM

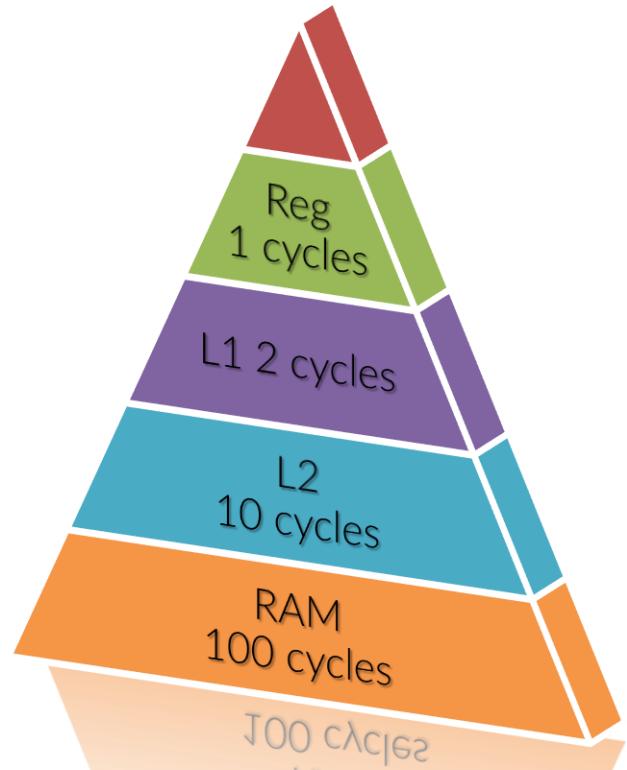
$$L_{1\text{cost}} + \text{Miss}_1 \times RAM_{\text{cost}}$$

- 100% L1 hit → 2 cycles
- 99% L1 hit → 3 cycles
- 97% L1 hit → 5 cycles

L1 cache + L2 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times (L2_{\text{cost}} + \text{Miss}_2 \times RAM_{\text{cost}})$$

- 100% L1 hit → 2 cycles
- 99% L1 hit, 100% L2 hit → 2.1cycles
- 97% L1 hit, 100% L2 hit → 2.3 cycles
- 90% L1 hit, 97% L2 hit → 3.3 cycles



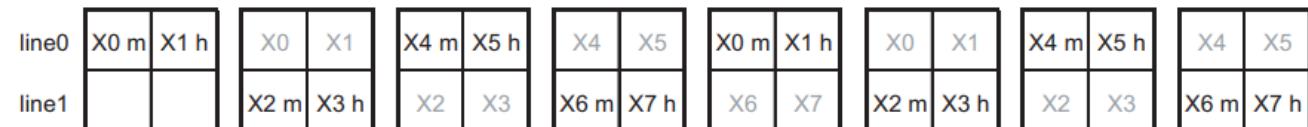


## When the cache is hit and when it is not: a simple model

Consider a simple direct mapped 16 byte data cache with two cache lines, each of size 8 bytes (two floats per line)

Consider the following code sequence, in which the array **X** is cache-aligned (that is, **X[0]** is always loaded into the beginning of the first cache line) and accessed twice in consecutive order:

```
float X[8];  
  
for(int j=0; j<2; j++)  
    for(int i=0; i<8; i++)  
        access(X[i]);
```



The hit-miss pattern is : MH MH MH MH MH MH MH MH,  
the miss-rate is 50% (the first miss is compulsory miss).



Let's consider another code sequence that access the array twice as before, but with a strided access

```
float X[8];
for(int j=0; j<2; j++)
{
    for(int i=0; i<7; i+=2)
        access(X[i]);
    for(int i=1; i<8; i+=2)
        access(X[i]);
}
```



The hit-miss pattern now is : MM MM MM MM MM MM MM MM,  
the miss-rate is 100%



Finally, consider a third code sequence that again access the array twice:

```
for(int i = 0; i < 2; i++)  
    for(int k = 0; k < 2; k++)  
        for(int j = 2*i; j < (i+1)*2; j ++)  
            access(X[ j ]);
```

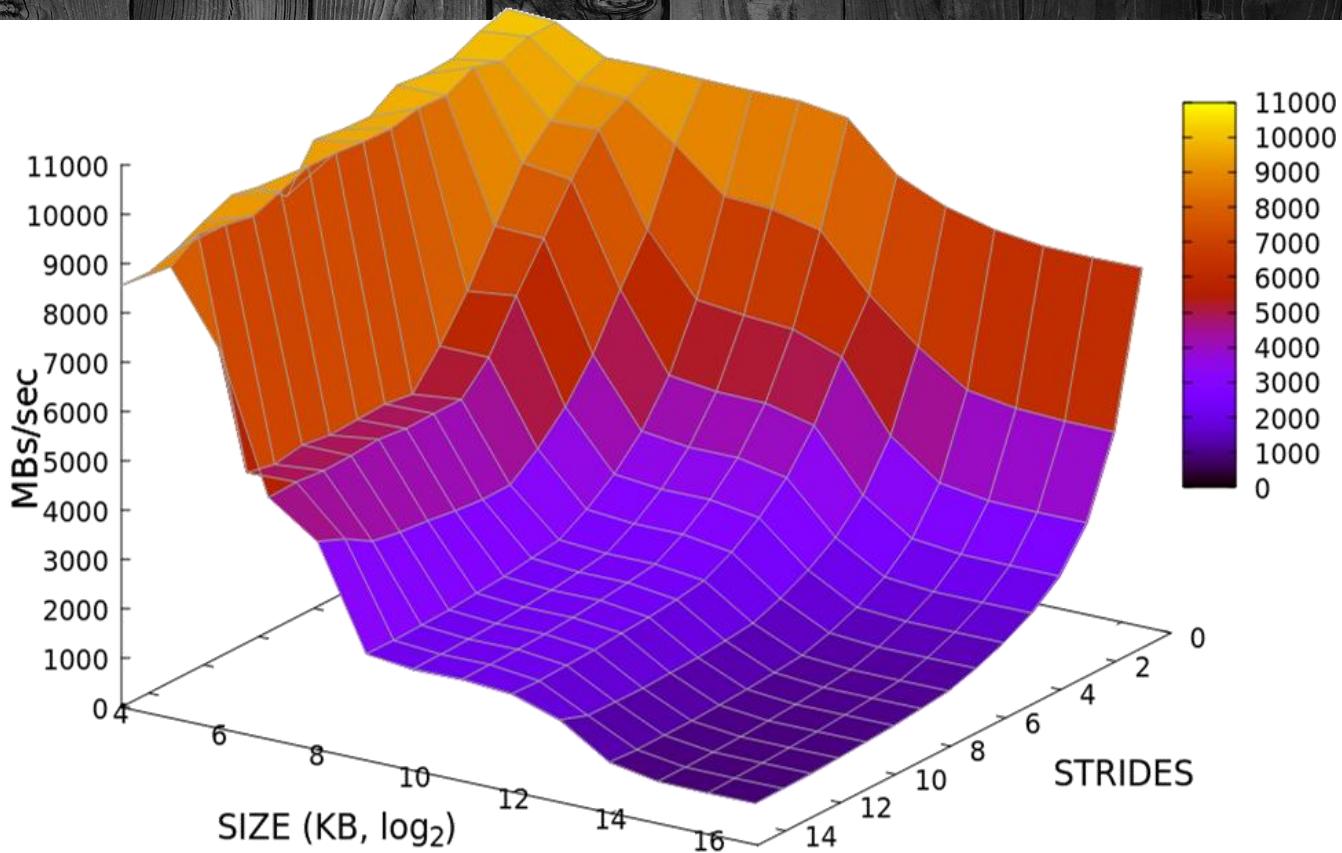


The hit-miss pattern now is : MH MH HH HH MH MH HH HH,  
the miss-rate is 25%

The main message is: memory access pattern is of primary importance



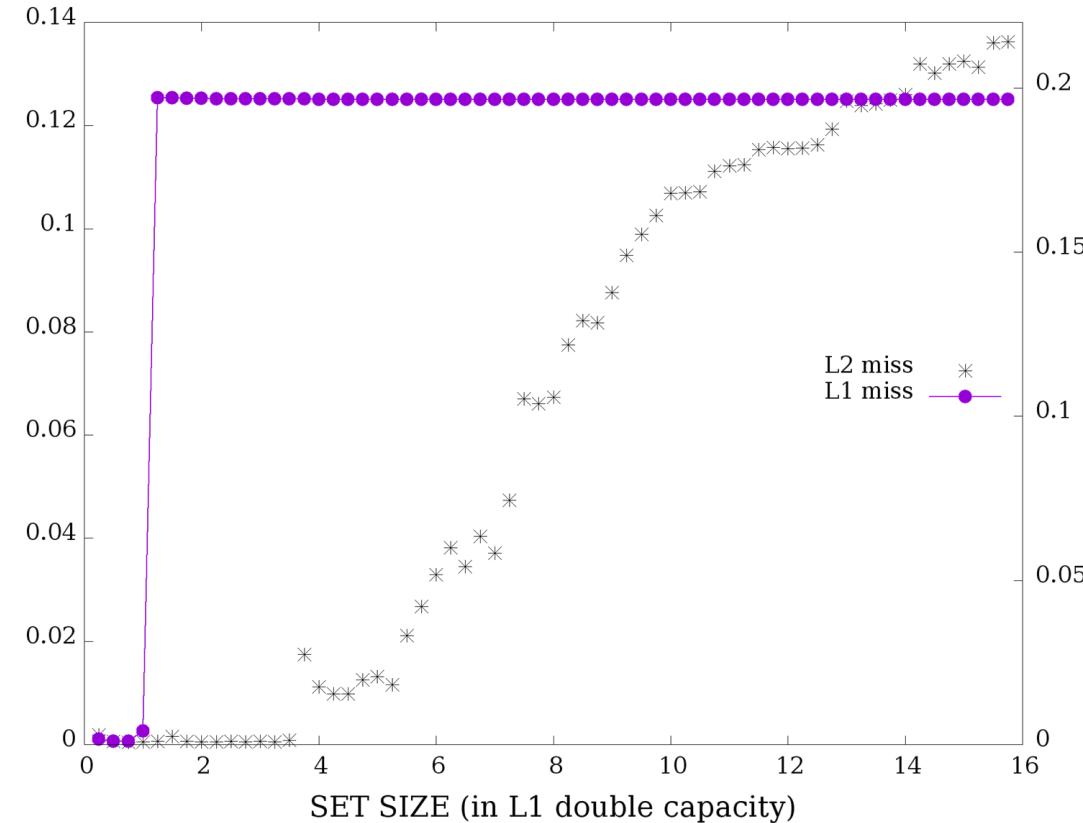
The result is..  
the memory mountain





Let's find out our cache size

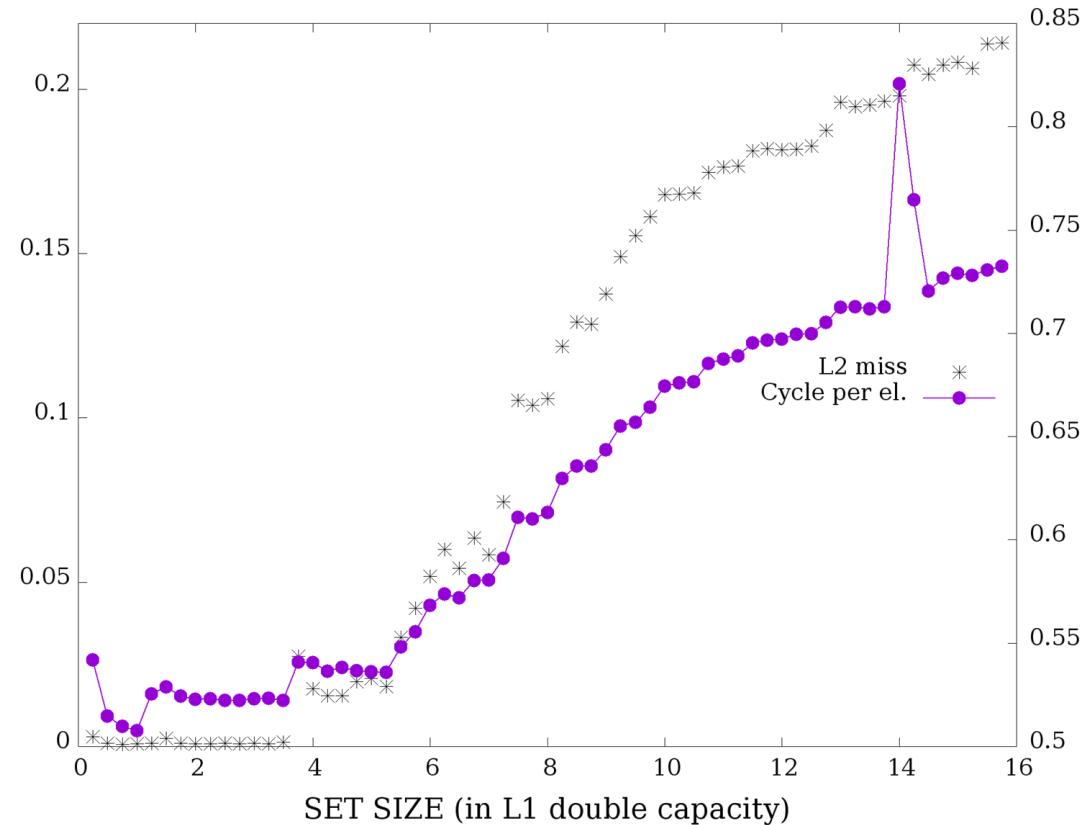
```
for (j=0; j<size; j++)  
    array[j] =  
        2.3*array[j]+1.2;
```





And the effect on  
cycles-per-operation  
metrics

```
for (j=0; j<size; j++)  
    array[j] =  
        2.3*array[j]+1.2;
```





## Matrix transpose

```
for(int row = 0; row < N; row++)  
    for(col = 0; col < N; col++)  
        A [ col*N + row ] = B [ row*N + col ];
```

NOTE: strided access to either A or B is unavoidable.

However: is it better to have it either on *read* or on *write* ?



Naïve version:

```
for(int row = 0; row < N; row++)  
    for(col = 0; col < N; col++)  
        A [ col*N + row ] = B [ row*N + col ];
```

Matrix transpose

NOTE: strided access to either A or B is unavoidable.

However: is it better to have it either on *read* or on *write* ?

Due to write-allocate transactions in the cache, **strided writes are more expensive than strided loads.**



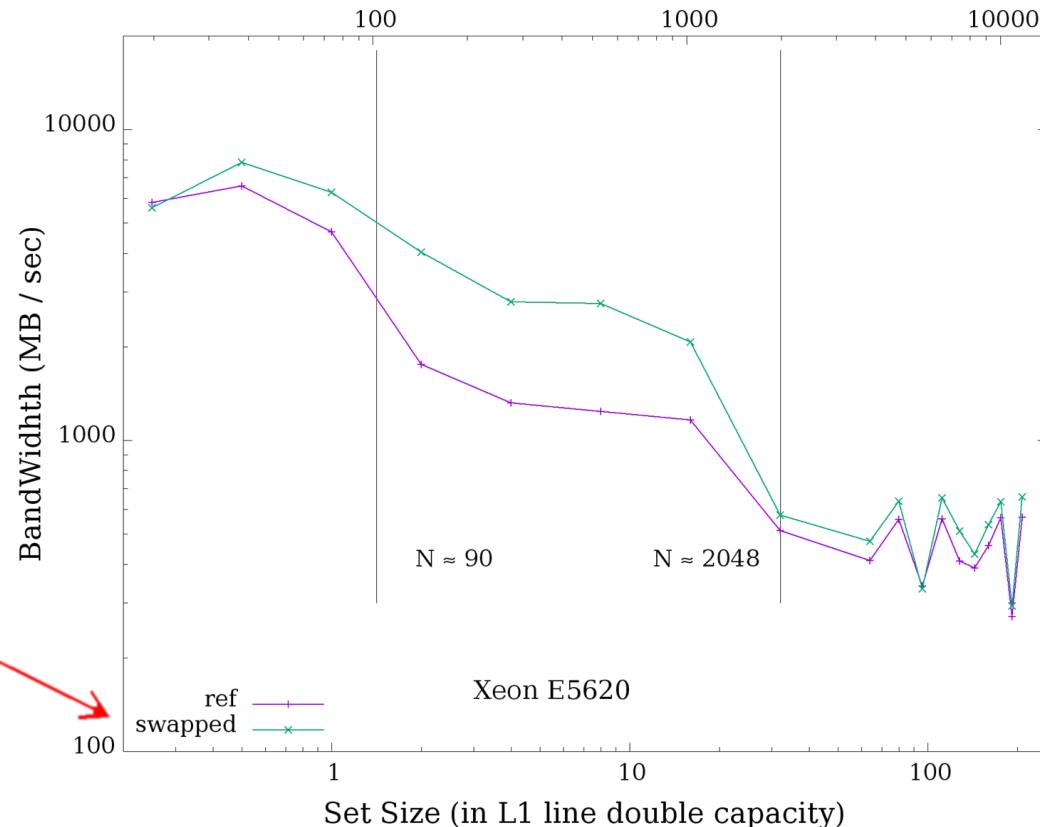
Inner cycle flipped:

$$A [ \text{row} * N + \text{col} ] =$$

$$B [ \text{col} * N + \text{row} ],$$



Due to time constraints, we can't go in deeper details.  
Ask if interested, though.





Reorder fields in structures so that what is used together stays together

### Linked-list node

```
struct my_node {  
    double   key;  
    char     my_data[300];  
    my_node *next_node; }
```

### Linked-list traversal

```
void myfunc(my_node *p, double key, <...>)  
{  
    while( p != NULL ) {  
        if( p->key == key ) {  
            do_something( <...> );  
            break;  
        }  
        p = p → next_node;  
    }  
}
```



Reorder fields in structures so that what is used together stays together

```
struct my_node
{
    double    key;
    char      my_data[300];
    my_node *next_node;
}
```



```
struct my_node
{
    double    key;
    my_node *next_node;
    char      my_data[300];
}
```



Split fields so that to keep consecutive the fields that are used sequentially

```
struct my_node
{
    double    key;
    char      my_data[300];
    my_node *next_node;
}
```

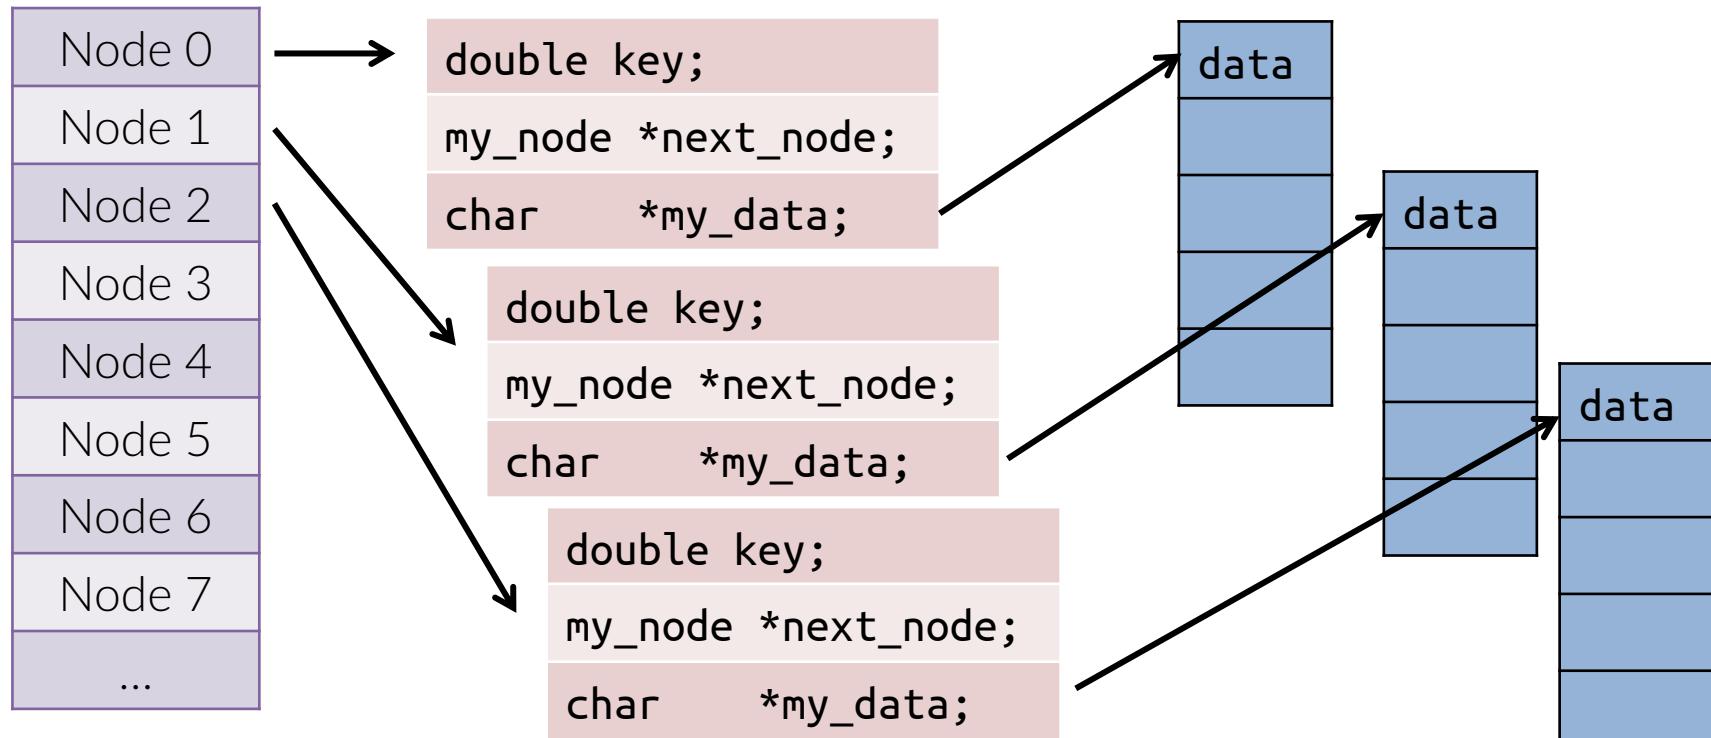


```
struct my_node
{
    double    key;
    my_node *next_node;
    void     *my_data;
}

struct my_data
{
    char data[300];
}
```

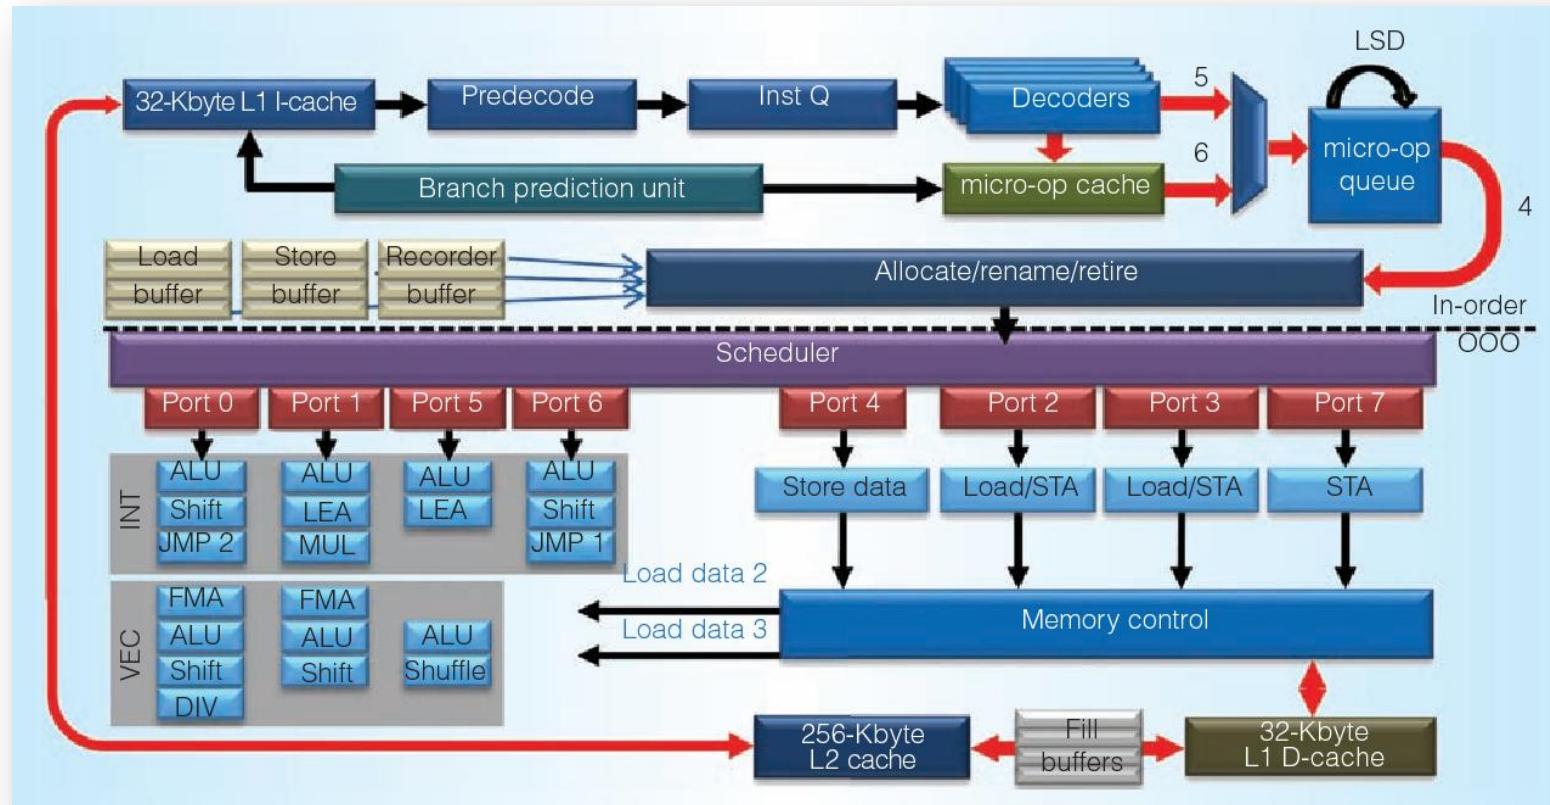


Those are called *hot* and *cold* fields



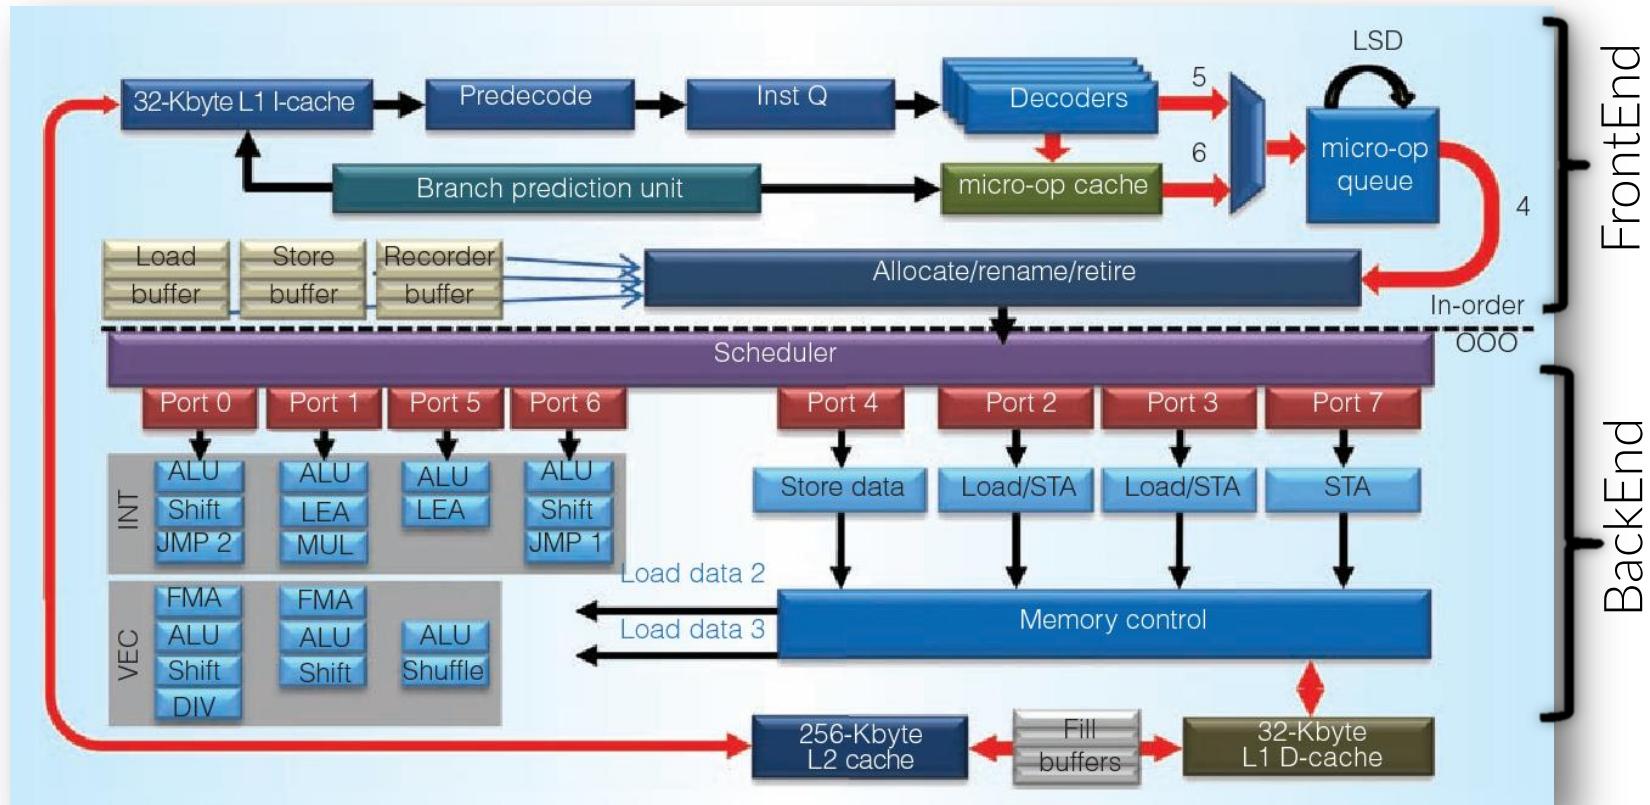


# Mid 90s: superscalar and out-of-order CPUs



6<sup>th</sup> generation  
SkyLake  
micro-arch.

More than 1 port is available to execute CPU instructions, although different units have different specializations (ALU, LEA, SHIFT, FMA, ... ) : that is superscalar capacity, i.e. the capacity of executing more than 1 instructions per cycle.



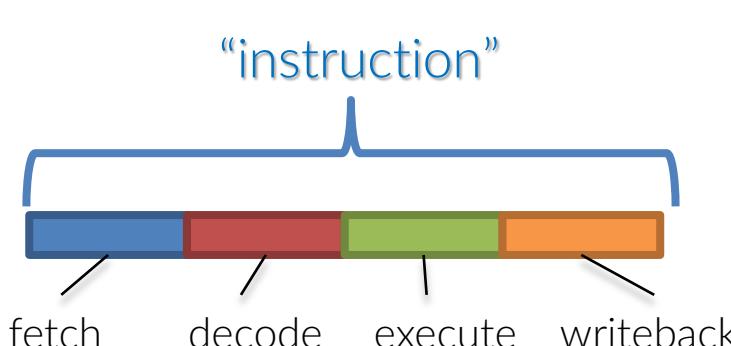
6<sup>th</sup> generation  
SkyLake  
micro-arch.



# Pipelines

It would be obvious to think that an “instruction” is a kind of *atomic* operation that the CPU perform as a whole. Indeed that was true until the mid of the 80s.

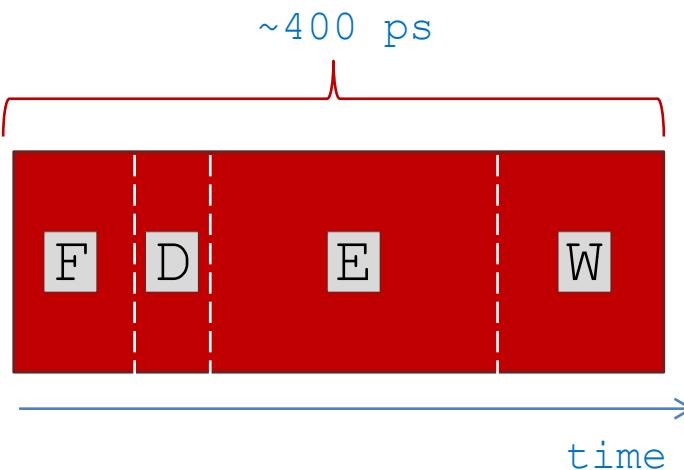
If you think carefully about it, it is easy to understand that actually an “instruction” involves at least the following *independent* steps:



1. ***Fetching***  
it must be recalled from memory/Icache
2. ***Decoding:***  
it must be “understood and interpreted”
3. ***Execution***
4. ***Writeback :***  
the result must be accounted in memory ()



# Pipelines



If all the four stages take ~400ps, we then would obtain a **throughput** of 2.5GIPS (giga-instructions per second).

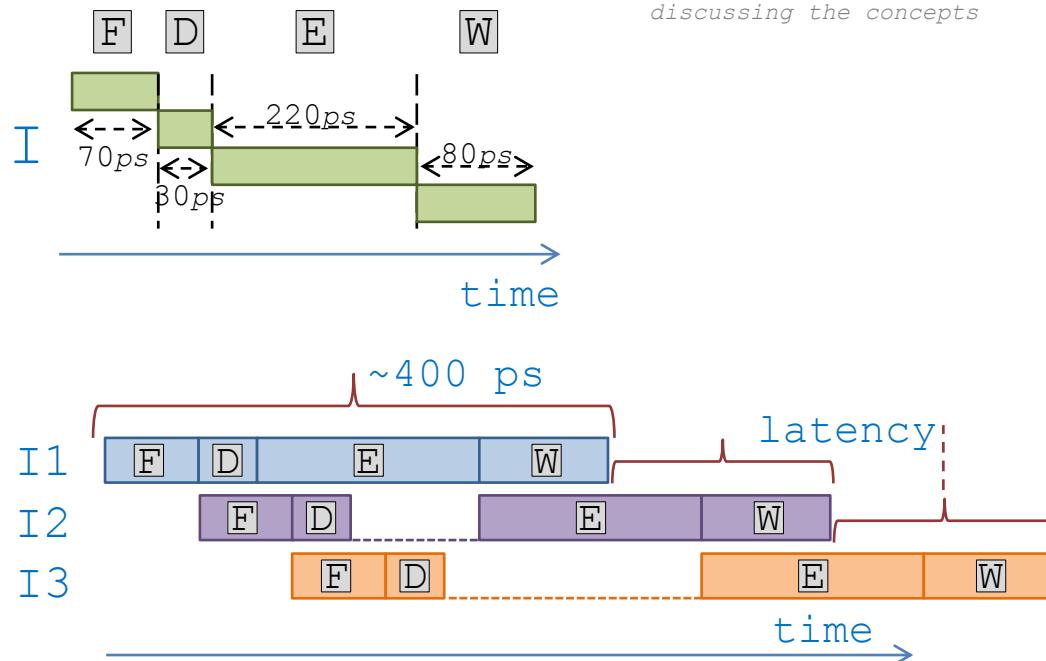
400ps is also the total time required to get a result from an instruction, and so it is the **latency** of the instruction.

However, if we were able to “detach” the four stages, we could organize things differently, like in a car-building chain, or even at the mensa of the university.



# Pipelines

Note: all the timing estimates are hypothetical for the purpose of discussing the concepts



If many independent logical units exist to perform each step, they could operate subsequently on different instructions:

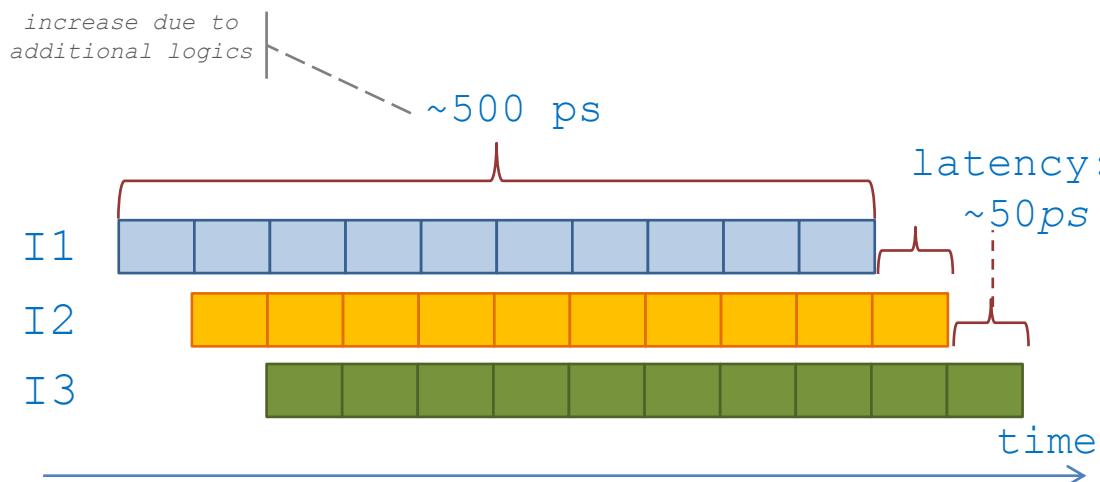
If the stage delays are not uniform, the throughput is limited by the latency  $F + (D+E) - (F+D) = E \sim 220\text{ps}$ , which means we have a throughput of  $\sim 4.5\text{GIPS}$  just because of logic units separation.



# Pipelines

Therefore, introducing the instructions pipelining, we can increase the **throughput** of our system by a large factor.

However, the efficiency of the pipelines is limited by its longest stage: the better option would be to have all equal stages, for instance further subdividing each stage – especially the most demanding ones (\*).



Now the throughput of our system has increased to 1 instruction retired every 50ps, i.e. 20GIPS

(\*) this is called *superpipelining*: modern CPUs may have 10-20 stages per pipeline.



# Pipelines

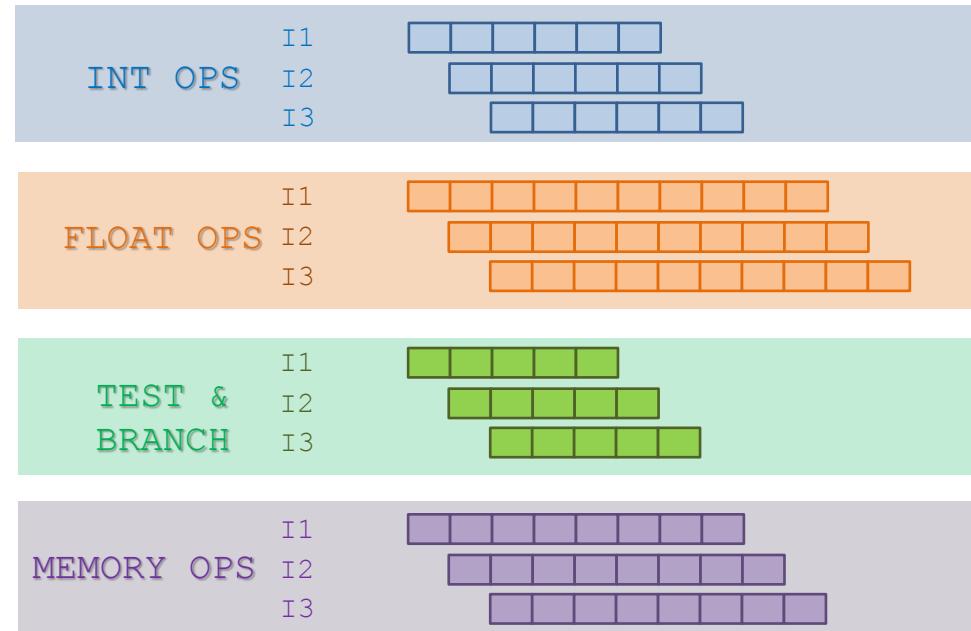
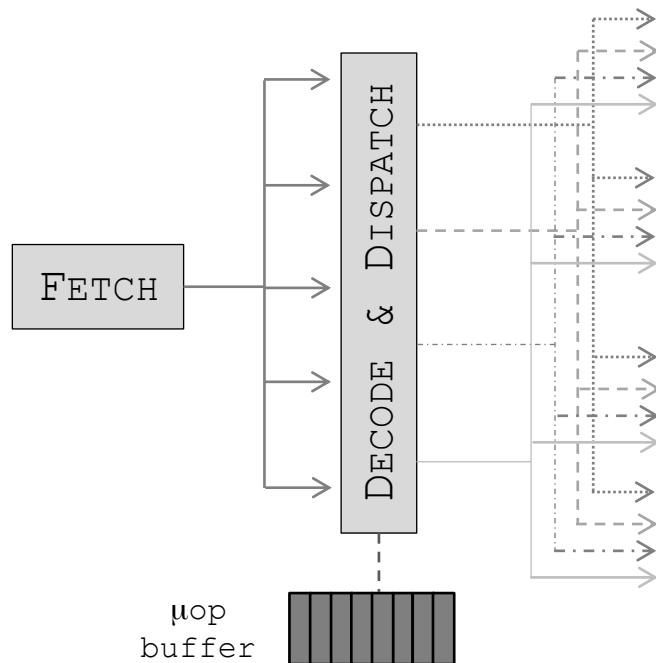
Pipelining is then about increasing the throughput of a system, and as we have seen it could be really effective.

However, there is more that can be done working on the **execution stage** that, of course, encompasses a large number of different *functional units*, performing a **different and independent tasks**.

With an enhancement of the decode/dispatch stage, we can address multiple pipelines that can be active at the same time:



# Pipelines





# Pipelines

The number of cycles between when an instruction's execution stage begins and when its result is available for other instructions usage, is the *latency* of the instruction and it grows with the pipeline deepness (the number of stages).

Typical latencies in modern processors (read the manual of yours) range from 1cyc for integer ops, to 3-6cyc for add and mul flop to  $\geq 10\text{-}20$  for div flop,  $\sim 50$  for trig flop.

Latency for memory loads can be severely troublesome, because they are highly unpredictable and they block all other operations making it difficult to fill the delay with some other ops.



# Pipelines hazards

Pipelining a system with *feedbacks* (i.e.: in which a result from an instruction can be the input of a following instruction, or can change the workflow) can expose the system to *hazards* when there are dependencies between successive instructions.

There are two forms of dependency:

1. *data dependency*: the result computed by the  $i^{th}$  instruction are used by one or more following instructions;
2. *control dependency*: the result of an instruction determines the value of the pc, i.e. the location of the next instruction to be executed.

When a dependency can lead to a wrong result, it is called an *hazard*.



# Pipelines hazards

If a control hazard arise, the entire pipeline content has to be flushed away, and a new instruction has to start from the beginning, then wasting a lot of CPU cycles.

This is one of the main cause of performance loss in modern superscalar-superpipeline-out-of-order CPUs. The logics for the runtime branch predictor occupies a large space on processor chips but it definitely is worth it.



# Pipelines hazards

Best branch predictors are as good as 95% of accuracy: nonetheless, the branch mis-prediction, or branch miss, determines a huge performance loss: typical values for penalty are 10-20 cycles!

That is because the longer the pipeline, the further in the future you have to scrutinize the flow, the more difficult it is and the larger will be the misprediction penalty.

## Example

Let's say we have 140 instructions on the flight, and 1 every 7 is a branching instructions. What is the probability that the pipelines shall not be flushed with 95% correct branch predictor? And with a 90% one?

~36% and ~12%



# Pipelines hazards

A very long pipeline, then, is not much more effective than a shorter one due to the real intrinsic nature of the codes that run on the CPUs.

The hazards we have just described are actually very common (very rarely a program is a stream of totally independent instructions with no jumps) and so the full exploitation of superscalarity + superpipelining is never reached.

Basically, that's the reason why we do not have 100-stage deep pipelines.

And the reason why branching is a performance-killer, too.



# Pipeline optimization

As we have seen, modern processor may be able to “perform more than one operations at a time”, through super-pipelining operations.

As for what concerns the programmer’s perspective, codes must be written so to make the pipelines saturation as effective as possible.

```
for (int i = 0; i < N; i++)  
    S += a[i] * b[i];
```

This way  $S$  is both read and written (the  $S$  as input in iteration  $i$  depends on  $S$  as output of iteration  $i-1$ ) and the FP pipeline is difficult to be exploited (iteration  $i$  must unavoidably wait for iteration  $i-1$  to end, which takes ~3-6 cycles at least)



# Pipeline optimization

Make things easier for the compiler and the CPU

```
for (i = 0; i < N; i += 2) {  
    sum0 += a[i] * b[i];  
    sum1 += a[i+1] * b[i+1]; }
```

Unrolling make it easier for the CPU to saturate the pipelines.

```
for (i = 0; i < N; i += 2) {  
    double tmp0 = a[i] * b[i];  
    double tmp1 = a[i+1] * b[i+1];  
    sum0 += tmp0;  
    sum1 += tmp1; }
```

Separate load and multiply from addition



# Pipeline optimization

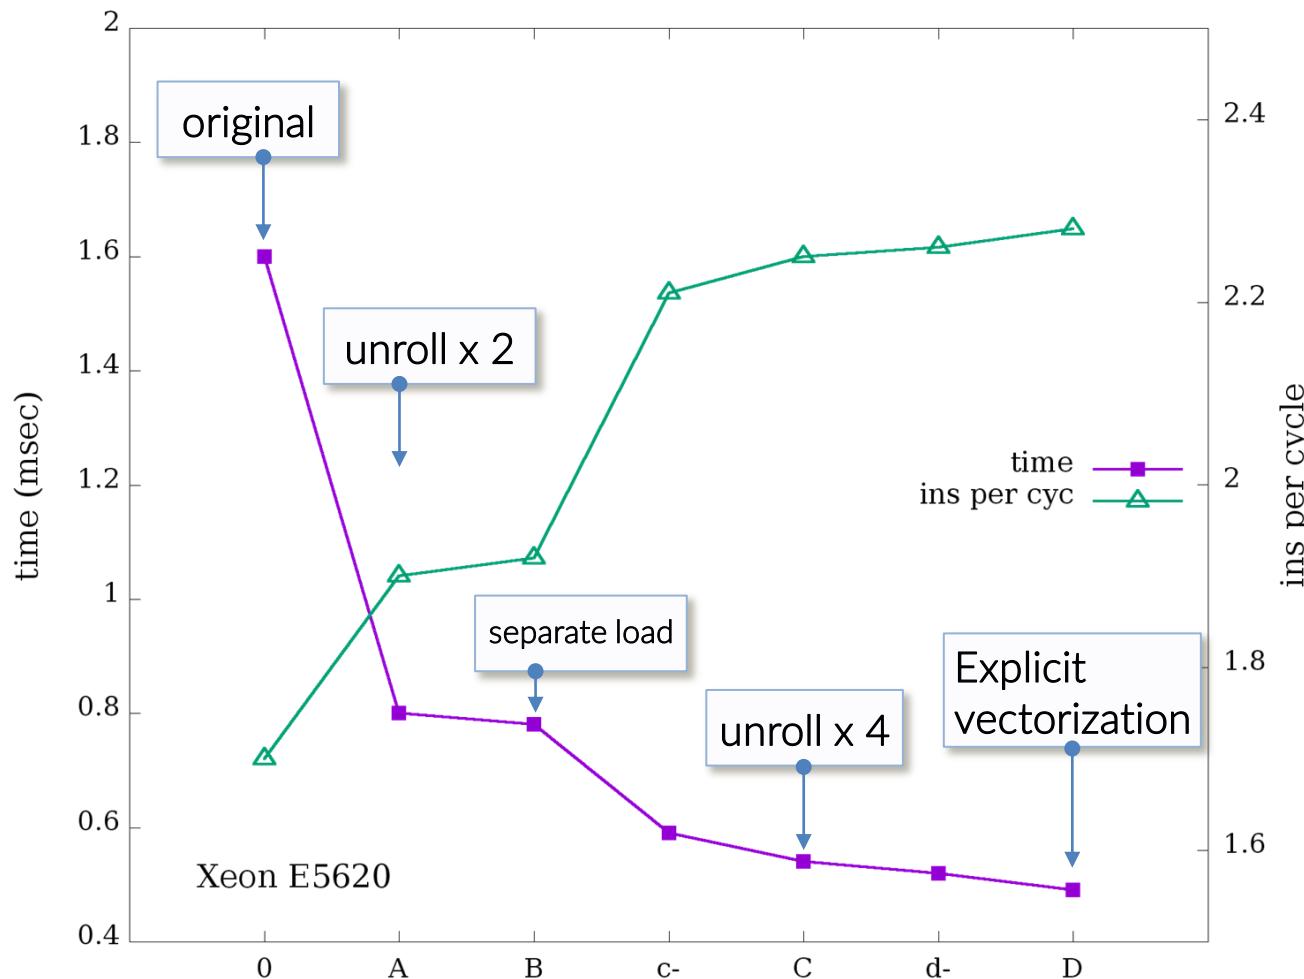
Make things easier for the compiler and the CPU

```
for (i = 0; i < N; i += 4) {  
    sum0 += a[i] * b[i];  
    sum1 += a[i+1] * b[i+1];  
    sum2 += a[i+2] * b[i+2];  
    sum3 += a[i+3] * b[i+3]; }
```

More unrolling may work even better

```
typedef double v4df __attribute__ ((vector_size (4*sizeof(double))));  
v4df array1, array2;  
v4df sum;  
for (i = 0; i < N/4; i++)  
    sum += array1[i] * array2[i];
```

Explicit vectorization





# Branch prediction

Conditional branches should be avoided as much as possible inside loops:

- moving them outside the loop and writing specialized loops
- performing variables/quantities set-up pree-emptively outside the loop
- using pointers to functions instead of selecting functions inside the loop
- substituting conditional branches with different operations



# Branch prediction

ex 1: Taking decisions before and outside the loop

```
for(i = 1; i < top; i++)
{
    if(case1 == 0) {
        if(case2 == 0) {
            if(case3 == 0)
                result += i;
            else
                result -= i;
        }
        else {
            if(case3 == 0)
                result *= i;
            else
                result /= i;
        }
    }
    else {
        if(case2 == 0) {
            if(case3 == 0)
                result += log10((double)i);
            else
                result -= log10((double)i);
        }
        else {
            if(case3 == 0)
                result *= sin((double)i);
            else
                result /= (sin((double)i) +
                           cos((double)i));
        }
    }
}
```



# Branch prediction

If you do not trust your compiler to perform the *loop hoisting* for you,

- define a specialized function for each case
- **before** and **outside** the loop set a function pointer to the right function

```
void (*func)(double *, int);
<here make func pointing to the right place>

double temp    = 0;
double result = 0;
for(i = 1; i < top; i++)
{
    func( & temp, i);
    result = temp;
}
```



# Branch prediction

However:

- Using function pointers you may incur in additional overhead due to function call.  
If the code snippets in different if-branches (or at least the most executed ones) are large/expensive, it might well be pointless (in modern CPUs).
- “Unrolling” the if-tree outside the for – then having multiple for loops may be highly unpractical if the branches are big piece of code.  
There’s no really a Swiss-knife recipe.

```
if (case1 == 0) {  
    if (case2 == 0) {  
        if (case3 == 0) {  
            for(i = 1; i < top; i++)  
                result += i;  
        } else  
            for(i = 1; i < top; i++)  
                result -= i;  
    }  
}
```



# Branch prediction

## ex 2: code restructuring

Consider the following code snippet

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```



# Branch prediction

Consider the following code snippet<sup>(\*)</sup>

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

qsort(data, SIZE, sizeof(int), compare);

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```

<sup>(\*)</sup>of course, you are adding an overhead due to the sorting routine, so the total running time may be even larger. Moreover, you should have all the values available so that does not work for real-time streamings. However, the point here is to focus on how – in general – it is better to avoid conditionals inside loop, with any possible trick or change in workflow



# Branch prediction

You can do even better, without adding operations:

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    t = (data[ii] - PIVOT -1) >> 31;
    sum += ~t & data[ii];
}
```



# Branch prediction

-00

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred
sum is 983597794767, elapsed seconds 5.40445
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow
sum is 983597794767, elapsed seconds 2.23186
(in total: 2.44473 seconds)
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart
sum is 983597794767, elapsed seconds 2.8878
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% █
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03
sum is 983597794767, elapsed seconds 0.660148
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03
sum is 983597794767, elapsed seconds: 0.650005
(in total: 0.795181 seconds)
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03
sum is 983597794767, elapsed seconds: 0.679286
```

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% █
```



# Branch prediction

-03

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03
sum is 983597794767, elapsed seconds: 0.660148
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03
sum is 983597794767, elapsed seconds: 0.650005
(in total: 0.795181 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03
sum is 983597794767, elapsed seconds: 0.679286
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% █
```

-03

-**march=native**

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03n
sum is 983597794767, elapsed seconds: 0.217864
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03n
sum is 983597794767, elapsed seconds: 0.215645
(in total: 0.355377 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03n
sum is 983597794767, elapsed seconds: 0.224288
```



# Branch prediction

What changes in the base version with -O3 ? → conditional move

Modern CPUs have the capability of performing *conditional move*, i.e to execute concurrently both branches of a conditional – if they are “simple enough” – and to select the right result upon the evaluation of the conditional

perform op1 → res in AX

perform op2 → res in BX

compare

if flag → mov BX in AX

HOWEVER: loops with conditionals can not be fully vectorized !!



# Branch prediction

Why the difference in the base v. between -O3 and -O3 -march=native ?

.L8:

```
movdqu  xmm0, XMMWORD PTR [rax]
movdqu  xmm6, XMMWORD PTR [rax]
movdqa  xmm2, xmm4
add     rax, 16
pcmpgt�  xmm0, xmm5
pand   xmm0, xmm6
pcmpgt�  xmm2, xmm0
movdqa  xmm3, xmm0
punpckldq  xmm3, xmm2
punpckhdq  xmm0, xmm2
paddq   xmm1, xmm3
paddq   xmm1, xmm0
cmp     rax, rcx
jne    .L8
```

compare **4 integers** at a time  
using xmmX registers, that are  
common to x86\_64  
architectures.

increase the counter by **4 int**

12 instructions to process **4 int**

.L8:

```
vmovdqu
add
vpcmpgt�
vpand
vpmovsxdq
vextracti128
vpaddq
vpmovsxdq
vpaddq
cmp
jne
.L8
```

bytes reshuffling  
to add one int at  
a time. These are SSE  
128-bits instructions

bytes reshuffling inside regs  
to add one int at a time  
The v prefix tells you these  
are AVX2 256-bits instr.

9 instructions to process **8 int**

compare **8 integers** at a time  
using ymmX registers. This  
requires AVX2 that is set on  
by -march=native for this  
CPU

increase the counter by **8 int**



# Branch prediction

We can do slightly better:

```
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```

can be changed to

```
for (ii = 0; ii < SIZE; ii++)
{
    acc = ( data[ii]>PIVOT )? data[ii] : 0;
    sum += acc;
}
```



# Branch prediction

Let's look to another practical example

You have 2 arrays, A and B, and you want to swap their elements so that

$$A[i] \geq B[i]$$

for all  $i$ .

A straightforward implementation would be:

```
for (i = 0; i < SIZE; i++)
{
    if ( A[i] < B[i] )
    {
        t = B[i];
        B[i] = A[i];
        A[i] = t;
    }
}
```



# Branch prediction

However, that implementation suffers exactly of the same problem we have just discussed.

An alternative way to write the same code, but in a more effective style is:

```
for (i = 0; i < SIZE; i++)
{
    int min = A[i] > B[i] ? B[i] : A[i];
    int max = A[i] >= B[i] ? A[i] : B[i];

    A[i] = max;
    B[i] = min;
}
```



# Branch prediction

```
for (uint ii = 0; ii < SIZE; ii++)  
{  
    if ( B[ii] > A[ii] )  
    {  
        int t = A[ii];  
        A[ii] = B[ii];  
        B[ii] = t;  
    }  
}
```

standard

```
for (uint ii = 0; ii < SIZE; ii++)  
{  
    int register t = -(A[ii]<B[ii]);  
    int register x = A[ii]^B[ii];  
    A[ii] = A[ii]^(x & t);  
    B[ii] = B[ii]^(x & t);  
}
```

smart2

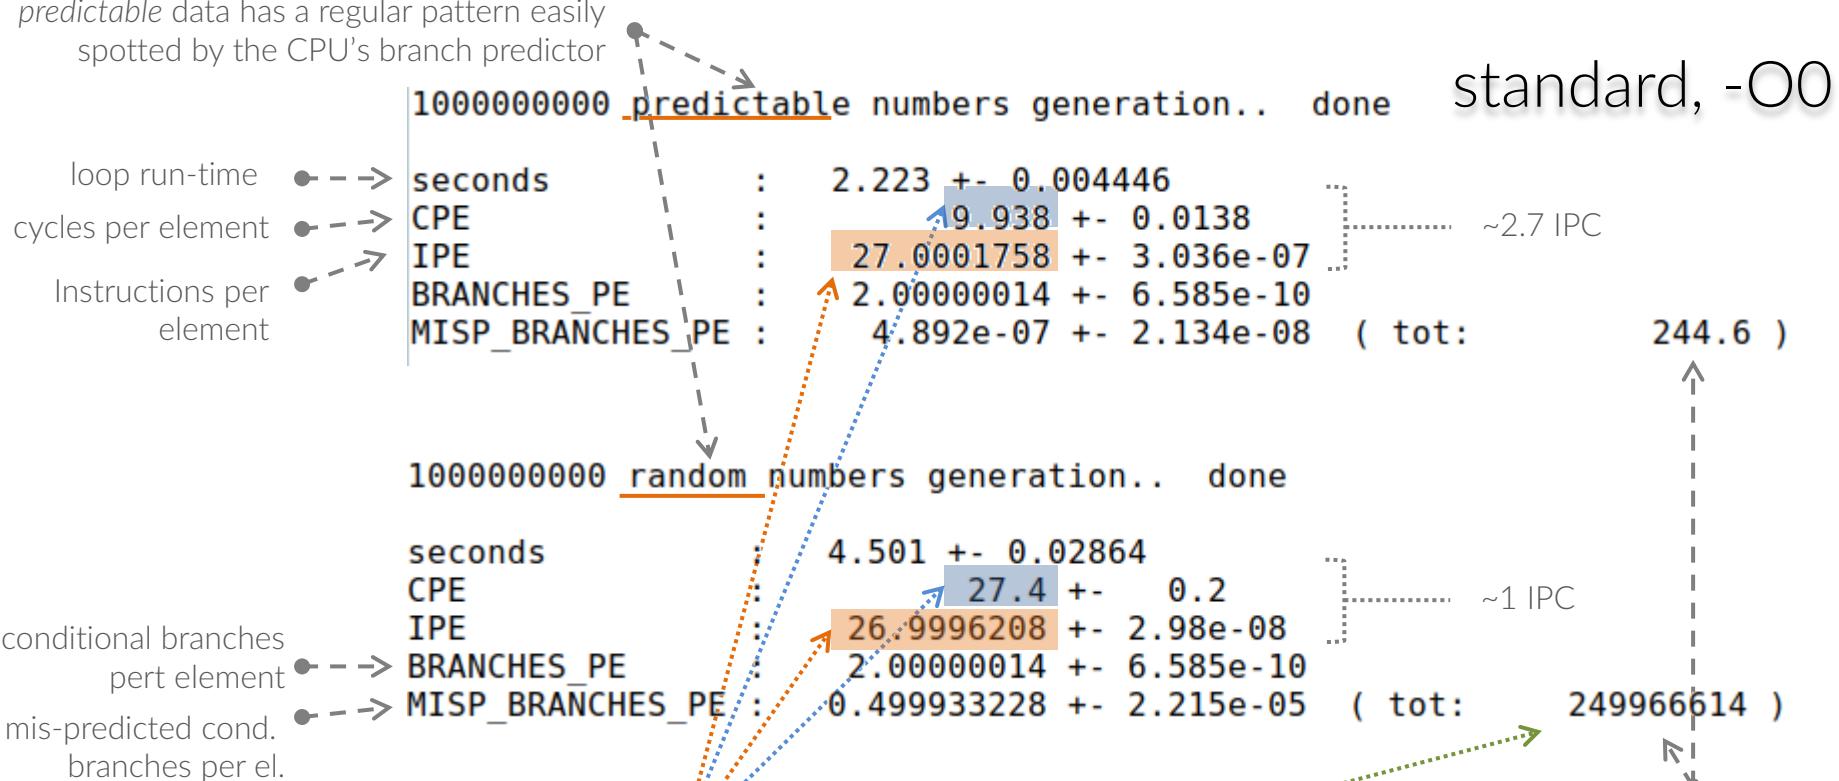
```
for (uint ii = 0; ii < SIZE; ii++)  
{  
    int max = (A[ii]>B[ii])? A[ii]:B[ii];  
    int min = (A[ii]>B[ii])? B[ii]:A[ii];  
    A[ii] = max;  
    B[ii] = min;  
}
```

smart

```
for (uint ii = 0; ii < SIZE; ii++)  
{  
    int d = A[ii]-B[ii];  
    d &= (d >> 31);  
    A[ii] = A[ii] - d;  
    B[ii] = B[ii] + d;  
}
```

smart3

*predictable* data has a regular pattern easily spotted by the CPU's branch predictor



same number of IPE but almost  
3 times as many CPE when  
data pattern is not predictable

in fact, there is 1 mis-predicted  
branch every 2  
( arrays are 500,000,000 long )

total # of mis-predicted  
conditional branches



# Branch prediction

seconds	1.824 +- 0.05734
CPE	11.12 +- 0.341
IPE	34.0000018 +- 4.344e-08
BRANCHES_PE	1.00000014 +- 3.293e-10
MISP_BRANCHES_PE	9.94e-07 +- 3.207e-08 ( tot: 497 )

number of IPE is larger than for standard case, but the CPE is stable !

**predictable**

mis-predicted branches are very few and comparable in both cases

seconds	1.857 +- 0.001762
CPE	11.32 +- 0.00192
IPE	34.0000018 +- 2.581e-08
BRANCHES_PE	1.00000014 +- 3.293e-10
MISP_BRANCHES_PE	9.512e-07 +- 4.987e-08 ( tot: 475.6 )

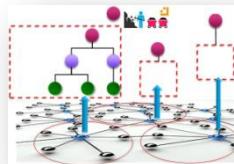
**random**



# A word of caution

It may be really easy to get lost in “optimization”, in hunting every single line wondering why some incredible trick that – you’re convinced – should work, actually does not.

“Optimization” includes also optimizing your effort and your time, so always remember that the **far most important ingredients** are:



The algorithms that you choose



The data model you design



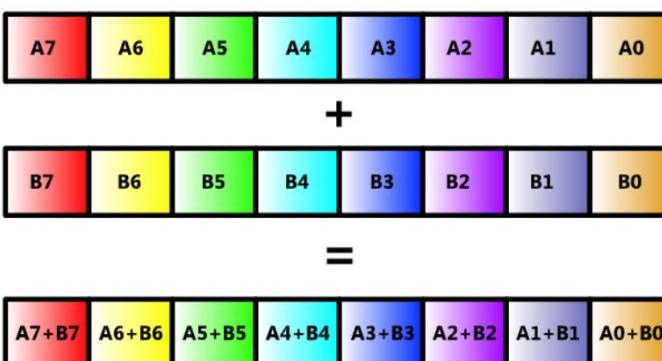
The overall quality, cleanliness and robustness of your code



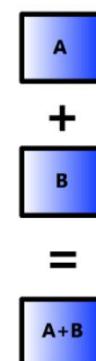
# Late 90s: vector capabilities become mainstream

[https://software.intel.com/sites/default/files/m/d/4/1/d/8/Intro\\_to\\_Intel\\_AVX.pdf](https://software.intel.com/sites/default/files/m/d/4/1/d/8/Intro_to_Intel_AVX.pdf)

## SIMD Mode



## Scalar Mode



Vector registers are large special registers in the CPU that can be considered subdivided in smaller independent chunks over which the same operation can be performed:

SIMD: Single Instruction  
Multiple Data



# Vector registers size

## SSE Data Types (16 XMM Registers)

<code>_m128</code>	Float	Float	Float	Float	4x 32-bit float
<code>_m128d</code>	Double		Double		2x 64-bit double
<code>_m128i</code>	B	B	B	B	16x 8-bit byte
<code>_m128i</code>	short	short	short	short	8x 16-bit short
<code>_m128i</code>	int		int	int	4x 32bit integer
<code>_m128i</code>	long long		long long		2x 64bit long
<code>_m128i</code>	doublequadword				1x 128-bit quad

## AVX Data Types (16 YMM Registers)

<code>_mm256</code>	Float	Float	Float	Float	Float	Float	Float	Float	8x 32-bit float
<code>_mm256d</code>	Double		Double		Double		Double		4x 64-bit double
<code>_mm256i</code>	<i>256-bit Integer registers. It behaves similarly to <code>_m128i</code>. Out of scope in AVX, useful on AVX2</i>								



# Early 00s: multi-core becomes mainstream

Late 00s : accelerators become mainstream

Early 10s: many-cores + heterogeneous computation  
become mainstream

That is the focus of the next two days...



# At the end.. first things first

“Optimization” may be a tag for several different concepts, as we have seen at the beginning of this course.

Many concurrent facts and factors must be kept together, and it is quite difficult to give general statements about which ones are more fundamental.

Top-level design plays a key role, as well as algorithm choice and implementation details.

Sometime – not at all that often! – even a single line of code, for instance a prefetching, may have brilliant consequences

...



# At the end.. first things first

However, unless you are seeking some extreme performance, as a general guideline just keep in mind that “optimization” reads

“let the compiler squeeze the maximum from your code”

Compilers are quite good indeed, and know the hardware they are running on better than you (99% of times).

So, as first, just learn how to :

- write non-obfuscated code
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures



# At the end.. first things first

- write non-obfuscated code
  - avoid memory aliasing
  - make it clear what a variable is used for and when
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures



# At the end.. first things first

- write non-obfuscated code
- design a good data structure layout
  - be cache-conscious
  - be NUMA-conscious
  - avoid race conditions in multi-threaded codes
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures



# At the end.. first things first

- write non-obfuscated code
- design a good data structure layout
- **design a “good” workflow**
  - compiler will be able to optimize branches and memory access patterns
  - prefetching will work better
  - make it easier to use multi-threading
- take advantage of the modern out-of-order, super-scalar, multi-core architectures



# At the end.. first things first

- write non-obfuscated code
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures
  - let the compiler exploit pipelining through operation ordering and unloop
  - let the compiler exploit the vectorization capabilities of CPUs
  - think task-based, data-driven



# C-specific hints

## Storage class

May load/store latency of heavily used variables

- **extern**  
Global variables, they exist forever
- **auto**  
Local variables, allocated on the stack for a limited scope, and then destroyed. They must be initialized
- **register**  
Suggests that the compiler puts this variable directly in a CPU register



# C-specific hints

## Variable qualifiers

- **const**  
Global variables, they exist forever
- **volatile**  
Indicates that this variable can be accessed from outside the program.
- **restrict**  
A memory address is accessed only via the specified pointer



# C-specific hints

:: note about **restrict** qualifier ::

```
void my_function( double *a, double *b, int n)
{
    for( int i = ; i < n; i++ )
        a[ i ] = s * b[ i - 1 ];
}
```

The compiler can not optimize the access to **a** and **b** because it can not assume that **a** and **b** are pointing to the same memory locations.

That is called *aliasing*, formally forbidden in fortran: which is the reason why in some cases fortran may compile in faster executables.

Help your C compiler in doing the best effort, either writing a clean code or using **restrict** or using **-fstrict-aliasing** **-Wstrict-aliasing** options.



# Compiler's usage

## Language standard

Some constructs can be better implemented if you specify a specific standard to the compiler.

## Architecture specification

The compiler is able to detect the architecture it runs on, but – to maximize executable's compatibility – it will not turn on very specific optimization unless you tell him so.

Read *carefully* the manual of your compiler.

In **gcc**, for instance, **-march=native** turns on a lot of sophisticated optimizations



# Compiler's usage

## Optimization level: -On

It is not granted that **-O3**, although often generating a faster code, is what you really need.

For instance, sometimes expensive optimizations may generate more code that on some architecture (e.g. with smaller Icaches) run slower, and using **-Os** may bring surprising results.

Take into accounts that modern compilers allows for local specific optimizations or compilation flags. In gcc for instance:

```
__attribute__ ((__option__ ("...")))  
__attribute__ ((optimize(n)))
```



# Compiler's usage

## Profile-guided optimization

Compilers (**gcc**, **icc** and **clang**) are able to instrument the code so to generate run-time information to be used in a subsequent compilations.

Knowing the typical execution patterns enables the compiler to perform more focused optimizations, especially if several branches are present.

For **gcc**:

Optimize branch prediction

```
gcc -fprofile-arcs
```

```
< ... RUN ... >
```

```
gcc -fbranch-  
probabilities
```

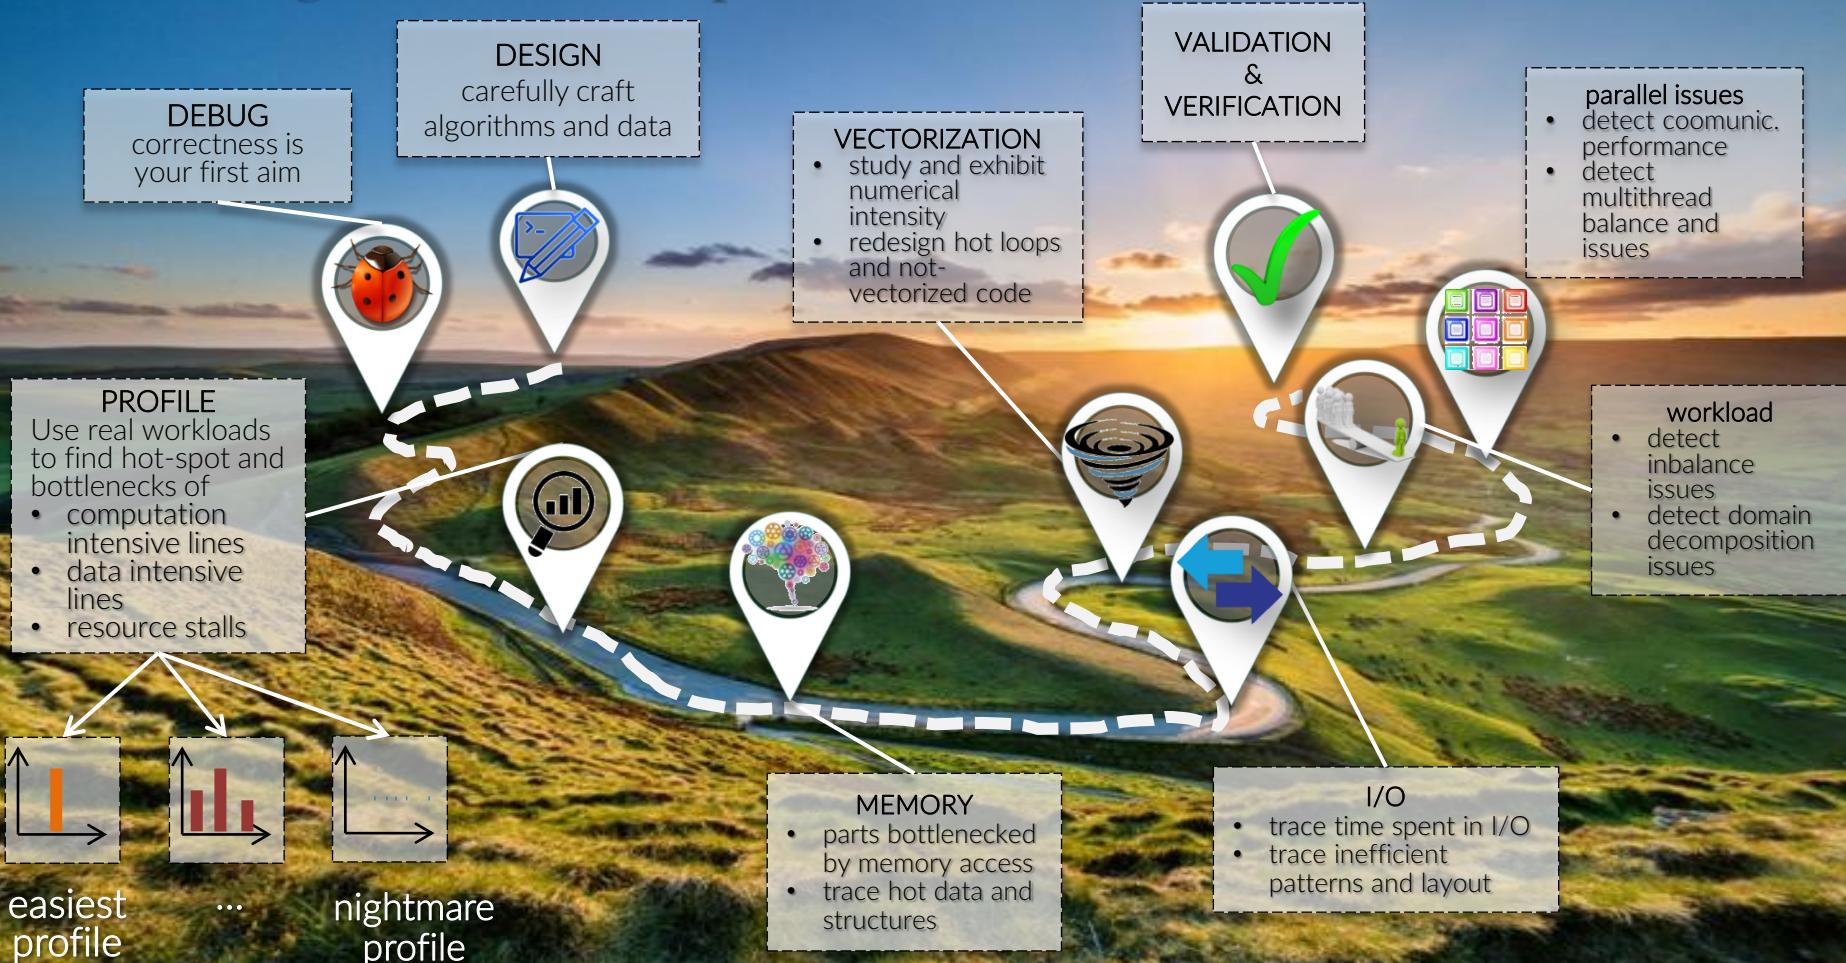
More general profiling

```
gcc -fprofile-generate
```

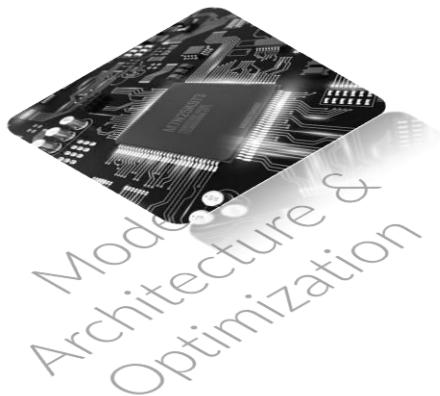
```
< ... RUN ... >
```

```
gcc -fprofile-use
```

# The winding road towards optimization



# Outline



Profiling



Debugging



# Map the workflow

A call tree (more precisely, a **call graph**) is a control flow graph that exhibits the calling relationships among routines in a program.

A *node* in the graph represent a routine, while an *edge* represents a calling relationship.

We'll concentrate on **dynamic call graphs** – i.e. the records of program executions.

The most complete graph is context-sensitive, which means that every call stack of a procedure is recorded as a separate node (the resulting graph is called *calling context tree* instead of *call tree*).

However, that requires a larger amount of memory for large program, it is useful in case of code reuse (the same code being executed at different points by different call paths).



# Basic profiling concepts

## INSTRUMENTATION

Inserts extra code at compile time wrapping function calls to count how many times it calls / is called and how much time it takes to execute.

## SAMPLING

The profiler ask for interrupts  $N_{samples}$  per sec + interrupts at function calls + interrupts at selected events, and records on a histogram the number of occurrences in every part of the program. The call graph is inferred from these data.

## DEBUGGING

The profiler ask for interrupts at every line code and function call (more correct: enters the by-step execution mode)



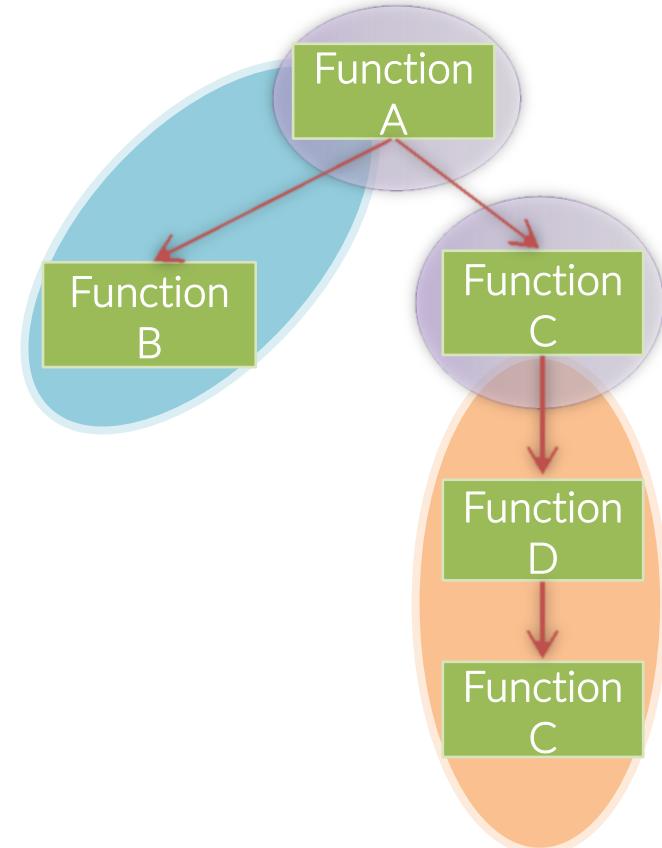
# Basic profiling concepts

## Collect program events

- Hardware interrupts
- Code instrumentation
- Instruction set simulation
- tracing

## Periodic sampling

- Top of the stack (exclusive)
- All stack (children inclusive)





# How to obtain the call tree

There are several way to obtain a dynamic call tree.

The main open source alternatives are:

1. Using `gcc` and `gprof`
2. Using linux `google perftools`
3. Using `valgrind`
4. Using `perf`
5. Some others, mostly non-free (among significant free: `CodeXL` by AMD)

Many IDE uses the aforementioned tools or their own plug-ins (`eclipse`, `netbeans`, `code::block`, `codelite`, ...)



## Using **gcc + gprof**

Just compile your source using **-pg** option.

You should also profile turning on the optimizations you're interested in.

```
gcc -[my optimizations] -pg -o myprogram.x myprogram.c
```

*Note: don't use the option **-p**. It provides less information than **-pg**.*

After that, run your program normally. Profiling infos will be written in the file **gmon.out**.

► You can read (options and details in the man page) the informations by  
**gprof myprogram.x**

► You can visualize the call graph by (read the man pages...)

```
gprof myprogram.x | gprof2dot.py | dot -T png -o callgraph.png
```



# The call-tree



Well, now that you just met, say goodbye to the glorious **gprof**

He's a dinosaur from the past decades..

- lacks real multithread support
- lacks real line-by-line capability
- does not profile shared libraries
- need recompilation
- may lie easier than other tools (see later..)

*You may still consider it for some call counts business*

*...may still consider it for some call counts business*

- *may lie easier than other tools (see later..)*
- *need recompilation*



# How to induce your profiler to lie

**gprof (gcc -pg)**  
does not record  
the call stack

```
#include <stdlib.h>
void loop(int n)
{
    int volatile i; // does not optimize out
    i = 0;
    while(i++ < n);
}

void light(int n) { loop(n); } void heavy(int n) {
loop(n); }

int main(void)
{
    light(100000);
    heavy(100000000);
    return 0;
}
```



## Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
101.30	0.20	0.20	2	101.30	101.30	loop
0.00	0.20	0.00	1	0.00	101.30	heavy
0.00	0.20	0.00	1	0.00	101.30	light



# How to induce your profiler to lie

How could it happen?  
The man page is clear:

We assume that the time for each execution of a function can be expressed by the total time for the function divided by the number of times the function is called.

Thus the time propagated along the call graph arcs to the function's parents is directly proportional to the number of times that arc is traversed.



# How to induce your profiler to lie

VALGRIND seems to better understand the situation, because it doesn't record only the number of calls to a function but also the time spent in a function given a call path (which, at odds, gprof infers backwardly).

That's why the valgrind measure of time spent in a function plus its callees is reliable.

However, it may as well end up to a misleading picture if we stick in an additional layer of complexity, unless you explicitly tell it to track separately different call stacks with the command-line option -

**-separate-callers = N**

Type	Callers	All Callers	Callee Map	Source Code
#	Ir	Source		
0		---		From '/home/luca/code/HPC_LECTURES/p'
14				void heavy(int n) { loop(n); }
15				
16		int main(void)		
17	0.00 {			
18	0.00 light(100000);			
19	0.10 1 call(s) to 'light' (lie1.g: lie1.c)			
19	0.00 heavy(100000000);			
19	99.88 1 call(s) to 'heavy' (lie1.g: lie1.c)			
20	0.00 return 0;			
21	0.00 }			
22				



## Using google perftools

That is the CPU profiler used at Google's. It provides a **thread-caching malloc**, a **heap-checker**, **heap-profiler**, **CPU profiler**.

As for the latter, basically there are 3 phases:

1. linking the library to the executable
2. running the executable
3. analyzing results

### [1] LINKING

There are two options:

- Link `-lprofiler` at the executable
- Adding the profiler at run time `LD_PRELOAD="/usr/lib/libprofiler.so"`  
`/path/to/exec`

This does not start the CPU profiling, though.



## Using google perftools

## [2] RUNNING

- Define

```
env CPUPROFILE= exec.prof /path/to/exec
```

You may define a signal, too

```
env CPUPROFILE= exec.prof /path/to/exec \
CPUPROFILESIGNAL= XX exec.prof &
```

so that to be able to trigger the start and stop of the profiling by

```
killall -XX exec
```

- In the code, include `<gperftools/profiler.h>` and encompass the code segment to be profiled within

```
ProfilerStart("name_of_profile_file")
...
ProfilerStop()
```

`CPUPROFILE_FREQUENCY=x` modifies the sampling frequency



## Using google perftools

## [3] ANALYZING RESULTS

```
pprof exec exec.prof
```

“interactive mode”

```
pprof --text exec exec.prof
```

output one line per procedure

```
pprof --gv exec exec.prof
```

annotated call graph via ‘gv’

```
pprof --gv --focus = some_func ...
```

restrict to code paths including “\*some\_func”

```
pprof --list = some_func ...
```

per-line annotated list of some\_func

```
pprof --disasm= some_func ...
```

annotated disassembly of some\_func

```
pprof -callgrind exec exec.prof
```

output call infos in callgrind format



## Using perf

### [1] COMPILING

```
gcc -g -fno-omit-frame-pointer my_prog.c -o my_prog
```

### [2] RUNNING

```
perf record -F ffff -call-graph <fp|lbr|dwarf> \
my_prog.c <args>
```

### [3] ANALYZING

```
perf report --call-graph=graph < --stdio >
```



## Using valgrind

### [1] COMPILE

```
gcc -g -fno-omit-frame-pointer my_prog.c -o my_prog
```

### [2] RUN

```
Valgrind --tool=callgrind --callgrind-out-file= $CALLGRIND_OUT --dump-instr=yes --coolect-jumps=yes --cache-sim=yes --branch-sim=yes < --I1=... >  
< --D1=...> my_prog <args>
```

### [3] ANALYZE

```
kcacheGrind $CALLGRIND_OUT
```



- BASIC / SYSTEM tools
  - gprof / gdb / perf / gperftools
  - Valgrind – cachegrind, callgrind,
  - ..
- HARDWARE COUNTER / PMU interface
  - perf
  - PAPI
  - Intel PMI
  - Likwid
  - ...
- HPC Tools
  - HPCToolkit
  - OpenSpeedShop
  - TAU
  - SCOREP +
- VENDOR tools
  - ARM-Allinea
  - CodeXL (AMD)
  - Intel tools
  - ...



# Tools : valgrind

An instrumentation framework for building dynamic analysis tools.

Valgrind basically runs your code in a virtual “sandbox” where a synthetic CPU (the same you have) is simulated and executes an instrumented code.

There are various Valgrind based tools for debugging and profiling purposes.

- Memcheck is a memory error detector → correctness
- Cachegrind is a cache and branch-prediction profiler → velocity
- Callgrind is a call-graph generating cache profiler. It has some overlap with Cachegrind
- Helgrind is a thread error detector → correctness
- DRD is also a thread error detector.  
Different analysis technique than Helgrind
- Massif is a heap profiler → memory efficiency using less memory
- DHAT is a different kind of heap profiler → memory layout inefficiencies
- SGcheck (experimental tool) that can detect overruns of stack and global arrays

KCacheGrind is a very useful GUI



## Memcheck : highlighting memory errors

- Invalid memory access: overrunning/underrunning of heap blocks or top of stack, addressing freed blocks, ...
- Use of variables with undefined values
- Incorrect freeing of heap memory
- Errors in moving memory (unwanted src/dst overlaps, ...)
- Memory leaks

## Cachegrind: simulating the cache

It can report how many hits (L1, L2 and L3, I- and D-) and how many misses.

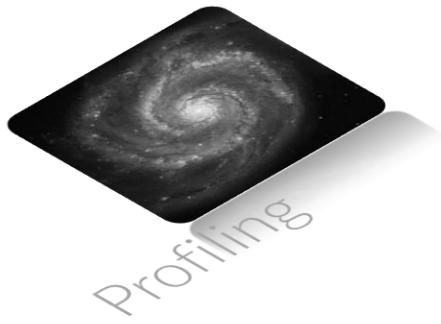
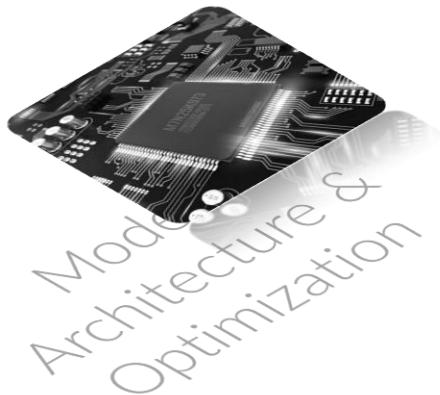
It can analyze CPU's branch prediction.

I<sub>r</sub>, I<sub>1mr</sub>, L<sub>1mr</sub>, D<sub>r</sub>, D<sub>1mr</sub>, D<sub>Lmr</sub>, B<sub>c</sub>, B<sub>cm</sub>, B<sub>i</sub>, B<sub>im</sub>, ...

## Callgrind: profiling the CPU

It collects the number of instructions executed, links them to source lines, records the caller/callee relationship between functions, and the numbers of such calls. It can collect data on cache simulation and/or branches.

# Outline



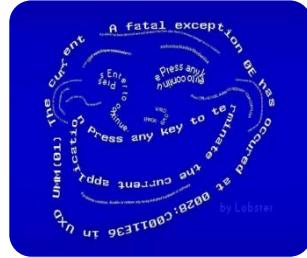
Debugging



# Debugging outline



```
if(Q > 0)
{
    switch(Q)
    {
        case(1): // row y = 0 and/or plane symmetry
            if(subregions[i][BOTTOM][x] == 0)
                // this subregion in quadrant 1 contains the
                // set-up corners for seed sub-region generation
                SBL[y] = 0, SBL[x] = Nmesh - subregion
                STR[y] = 1, STR[x] = Nmesh - subregion
                // find horizontal extension in this quadrant
                Np = STR[x] - SBL[x];
                // allocate memory for seeds in this string
                SEED_y = (unsigned int*)malloc(sizeof(unsigned
                int)*Np);
                if(!internal.nmic_original_seedtable)
                {
                    for(j = 0; j < Np; j++)
                        SEED_y[j] = internal.nmic_original_seedtable[j];
                }
            }
        }
}
```



Introduction  
Debugging  
a code

Inspecting  
a code crash

Debugging  
a running  
process



# Introducing debugging



## Smithsonian



It has long been recognized and documented that insects are the most diverse group of organisms, meaning that the numbers of species of insects are more than any other group. In the world, some 900 thousand different kinds of living insects are known. This representation approximates 80 percent of the world's species. The true figure of living species of insects can only be estimated from present and past studies. Most authorities agree that there are more insect species that have not been described (named by science) than there are insect species that have been previously named. Conservative estimates suggest that this figure is 2 million, but estimates extend to 30 million. In the last decade, much attention has been given to the entomofauna that exists in the canopies of tropical forests of the world. From studies conducted by Terry Erwin of the Smithsonian Institution's Department of Entomology in Latin American forest canopies, the number of living species of insects has been estimated to be 30 million. Insects also probably have the largest biomass of the terrestrial animals. At any time, it is estimated that there are some 10 quintillion (10,000,000,000,000,000,000) individual insects alive.



# Introducing debugging



9/9	
0 800	Autan started
1 000	stopped - autan ✓ { 1.2700 9.037 847 021 13° UC (03) MP - MC 1.362170000 9.037 846 995 (03) PRO 2 2.130476415 4.6159250 const 2.130676315
	Relays 6-2 in 033 failed special speed test in relay " 10.000 test . Relays changed
1 100	Started Cosine Tape (Sine check)
1 225	Started Multi Adder Test.
1 545	
1 600	First actual case of bug being found. Autang started.
1 700	closed down .

The usage of **bug** and, hence, of **de-bugging** referred to programming is a long-standing tradition, whose origin is difficult to trace back.

An often-told story is about Mark-II calculator located at Harvard: On Sept. 9<sup>th</sup>, 1945, a technician found a moth in a relay that caused a flaw in a “program” execution.

Adm. Grace Hopper is reported to have written in its diary about that as “first actual case of bug being found”

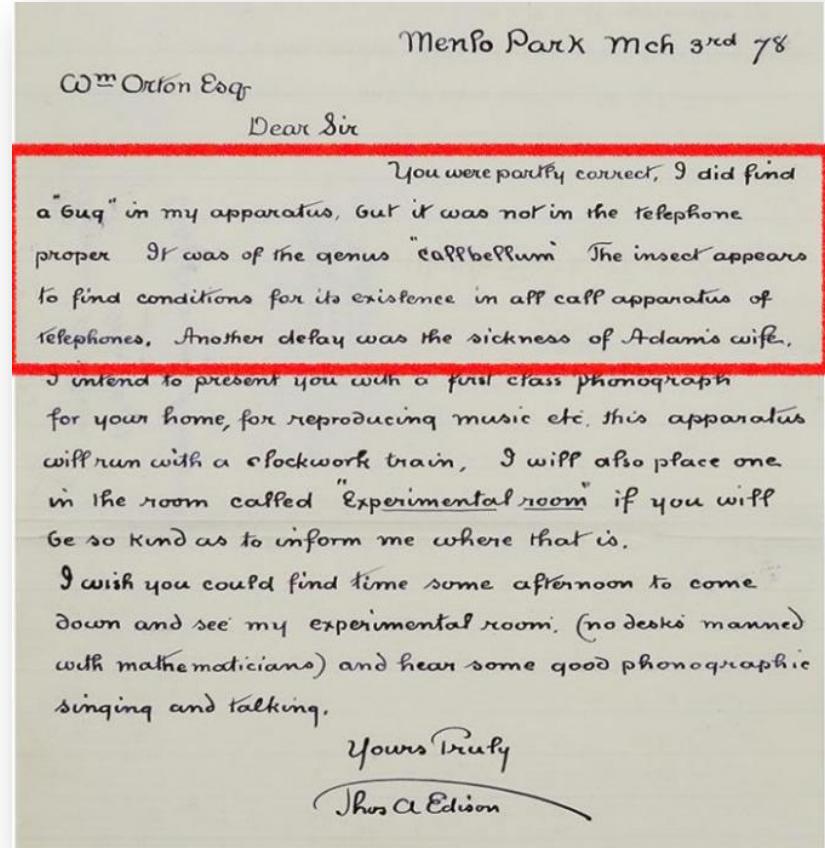


# Introducing debugging

However, Thomas Edison found another *actual* bug in one of his phones, as he reports in a letter to an associate.

He later writes:

"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that "**Bugs**"—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."





# Introducing debugging



It is difficult to trace back how this term has been poured into computer programming jargon.

However, already in the early 60s it was appearing in technical papers without need of explanation.

The immortal I. Asimov used it in a 1944 short robot story “Catch that rabbit”, and his incredible influence contributed much to make the term popular.

More funny infos:

[IEEE Annals of the History of Computing](#)  
( Volume: 20 , [Issue: 4](#) , Oct-Dec 1998 )

ask for the PDF if you're interested and do not have access

## **Stalking the Elusive Computer Bug**

PEGGY ALDRICH KIDWELL

*From at least the time of Thomas Edison, U.S. engineers have used the word “bug” to refer to flaws in the systems they developed. This short word conveniently covered a multitude of possible problems. It also suggested that difficulties were small and could be easily corrected. IBM engineers who installed the ASCC Mark I at Harvard University in 1944 taught the phrase to the staff there. Grace Murray Hopper used the word with particular enthusiasm in documents relating to her work. In 1947, when technicians building the Mark II computer at Harvard discovered a moth in one of the relays, they saved it as the first actual case of a bug being found. In the early 1950s, the terms “bug” and “debug,” as applied to computers and computer programs, began to appear not only in computer documentation but even in the popular press.*

### **Introduction**

**S**talking computer bugs—that is to say, finding errors in computer hardware and software—occupies and has occupied much of the time and ingenuity of the people who design, build, program, and use computers.<sup>1</sup> Early programmers realized this with some distress. Maurice Wilkes recalls that in about June of 1949:

I was trying to get working my first non-trivial program, which was one for the numerical integration of Airy's dif-

published in the *Annals* in 1981. Here the time is given as the summer of 1945, but the computer is the Mark II, not the Mark I. A photograph shows the moth taped in the logbook, labeled “first actual case of bug being found.” Small problems with computers have been called bugs ever since.<sup>5</sup>



# Introducing debugging

That maybe is too dramatic and emphatic, but the point to get from Edison is that the de-bugging activity is an inherent and intrinsic one in software development

"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that "**Bugs**"—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."

T. Edison



# Introducing debugging



Provided that

- you'll have close encounters with bugs in your life;
- de-bugging is a fundamental and unavoidable part of your work;

what is the best way to proceed ?



# Introducing debugging



## 1. Do not insert bugs

*Highly encouraged, but rarely works*



# Introducing debugging



## 1. Do not insert bugs

*Highly encouraged, but rarely works*

## 2. Add `printf` statements everywhere

*Highly discouraged, but sometimes works.*

*In case of memory problems – typically due to pointers chaos – you may see that the problem “disappear”, or changes its appearance, when you insert a new `printf`*



# Introducing debugging



## 1. Do not insert bugs

*Highly encouraged, but rarely works*

## 2. Add `printf` statements everywhere

*Highly discouraged, but sometimes works.*

*In case of memory problems – typically due to a os*

## 3. Use a DEBUGGER

*That is definitely the best choice and, fortunately,  
the subject of this lecture.*



# Introducing debugging



## 1. Do not insert bugs

*Highly recommended, but much harder.*

**gdb** is almost certainly the best free, extremely feature-rich command-line debugger ubiquitously available on \*nix systems.

## 3. Use a DEBUGGER

*That is definitely the best choice and, fortunately, the subject of this lecture.*



There are 3 basic usages<sup>(\*)</sup> of GDB:

1. Debugging a code  
*best if it has been compiled with `-g`*
2. Inspecting a code crash through a core file
3. Debugging / inspecting a running code

(\*) Highly advised: learn keyboard commands. Although many gui exist (we'll see some later), keyboard is still the best productivity tool.



# Debugging outline



Introduction

```
if(Q > 0)
{
    switch(Q)
    [
        case(1): // row y = 0 and/or plane symmetry
            if(subregions[i][BOTTOM][x] == 0)
                // this subregion in Quadrant 1 contains the
                // set-up corners for seed sub-region generation
                SBL[y] = 0, SBL[x] = Nmesh - subregion
                STR[y] = 1, STR[x] = Nmesh - subregion
                // Find horizontal extension in this quadrant
                NP = STR[x] - SBL[x];
                // allocate memory for seeds in this string
                SEED_y = (unsigned int*)malloc(sizeof(unsigned
                int)*NP);
                if(!internal.nmic_original_seedtable)
                {
                    for(j = 0; j < NP; j++)
                        SEED_y[j] = internal.nmic_original_seedtable[j];
                }
            }
        }
}
```



Debugging  
a code

Inspecting  
a code crash

Debugging  
a running  
process



# Compiling with dbg infos



In order to include debugging information in your code, you need to compile it with `-g` family options (read the gcc manual for complete info):

- g** produce dbg info in O.S. native format
- ggdb** produce gdb specific extended info, as much as possible
- glevel** default level is 2. 0 amounts to no info, 1 is minimal, 3 includes extra information (for instance, macros expansion) – this allows macro expansion; add **-gdwarf-n** in case, where possibile, where *n* is the maximum allowed (4)
- ggdblevel** you can combine the two to maximize the amount of useful info generated
  - remember: **-fno-omit-frame-pointer**, especially if you are using –Ox



Debugging  
a code

# Debugging: start



You just start your code under gdb control:

```
%> gdb program
```

You can define the arguments needed by your program already at invocation:

► live demo with `gdb_try_args.c`

```
%> gdb --args program arg1 arg2 ... argN
```

Or you can define the arguments from within the gdb session:

```
%> gdb program
Reading symbols from program...done.
(gdb) set args arg1 arg2 ... argN
(gdb) run
```



# Debugging: run & stop



You may just want the code to run, for instance to reach the point of a seg fault:

```
%> (gdb) run
```

Or, you may want to stop it from the beginning to have full control of each step:

```
%> (gdb) break main  
%> (gdb) run
```

Or you may already know what is the problematic point to stop at:

```
%> (gdb) break location  
%> (gdb) run
```



Breakpoints are a key concept in debugging. They are stopping point at which the execution interrupts and the control is given back to you, so that you can inspect the memory contents (variables values, registers values, ... ) or follow the subsequent execution step by step.

You can define a breakpoint in several way

(gdb) break

(gdb) break *±offset*

(gdb) break  
*filename:linenum*

(gdb) break  
*functionname*

There are more  
options, just check  
the manual

insert a break at the current pos

insert a break *offset* lines after/before the current line

insert a break at *linenum* of file *filename*

insert a break at the entry point of function *functionname*



# Debugging: run & stop



A breakpoint may be defined as dependent on a given condition:

```
(gdb) break my_function if (arg1 > 3 )
```

This sets a breakpoint at function *my\_function* : the condition will be evaluated each time the point is reached, and the execution is stopped only if it is true.

Condition can be any valid expression.

```
(gdb) info break
```

gives you informations on active breakpoints.

```
(gdb) delete [n]
      clear [location]
<disable | enable>  see the manual
```



You can define a list of commands to be executed when a given breakpoint is reached:

```
(gdb) break my_function if (arg1 > 3 )
Breakpoint 1 at 0x.....: file blabla.c, line 42
(gdb) command 1
Type commands for when breakpoint 1 is hit, one per
line.
```

End with a line saying just "end".

```
> print arg1
> print another_useful_variable
> x/10wd a_global_integer_array
```



# Debugging: run & stop



When you have the control of the program execution, you can decide how to proceed:

(gdb) cont [c]	continue until the end / next stop
(gdb) cont <i>count-ignore</i>	continue ignoring the next <i>count-ignore</i> stops (for instance, occurrences of a bp)
(gdb) next [n]   <i>count</i> nexti	continue to the next src line <i>in the current stack frame</i>
(gdb) step [s]   <i>count</i> stepi	continue to the next src line
(gdb) until [u]   <i>count</i>	continue until a src line past the current one is reached <i>in the current stack fr.</i>
(gdb) advance <i>location</i>	continue until the specified location is reached

► live demo with `gdb_try_breaks.c`



Debugging  
a code

# Debugging: run & stop



You can also rewind your execution step-by-step

```
(gdb) reverse-continue [rc]  
  
(gdb) reverse-step [count]  
reverse-stepi  
  
(gdb) reverse-next [count]  
reverse-nexti  
  
(gdb) set exec-direction <verse | forward >
```



# Debugging: source list



Often, when you are debugging, you may have the need of looking at either the source lines or at the generated assembler:

(gdb) *list linenum*

print src lines around line *linenum* of the current source file

(gdb) *list function*

print the source lines of *function*

(gdb) *list location*

print src lines around *location*

(gdb) *set listsize count*

control the number of src lines printed

(gdb) *disass*  
[*/m*] [*function*] [*location*]

show the assembler

► live demo with `gdb_try_breaks.c`



Examining the stack is often of vital importance. With GDB you can have a quick and detailed inspection of all the stack frames.

(gdb) `backtrace [args]`

*n*

`-n`

`full`

print the backtrace of the whole stack

print only the *n* innermost frames

print only the *n* outermost frames

print local variables value, also

`where, info stack`

additional aliases

► live demo with `gdb_try_breaks.c`



Accessing to the content of memory is a fundamental ability of a debugger. You have several different ways to do that:

(gdb) *print variable*  
p/F *variable*

print the value of *variable*  
print *variable* in a different format  
(x, d, u, o, t, a, c, f, s)  
see then manual for advanced location

(gdb) x/FMT *address*

Explore memory starting at address *address*  
-> see at live demo how to use this

(gdb) *display expr*  
display/fmt *expr*  
display/fmt *addr*

Add *expr* to the list of expressions to display each time your program stops

► live demo with `gdb_try_breaks.c`



Debugging  
a code

# Debugging: memory examination



Memory can be searched to find a particular value, of a given size

► live demo with `gdb_try_breaks.c`

(gdb) `find [/sn] start,  
end, val1 [,val2, ...]`

`find [/sn] start,  
+len, val1 [,val2, ...]`

Search memory for a particular sequence of bytes.  
s is the size of type to be searched  
n is the max number of occurrences



Debugging  
a code

# Debugging: registers examination



Examining registers may be also useful (although it's something that only quite advanced users can conceive)

- (gdb) info registers print the value of all registers
- (gdb) info vector print the content of vector registers
- (gdb) print \$rsp print value of the stack pointer
- (gdb) x/10wd \$rsp print values of the first 10 4-bytes integers on the stack
- (gdb) x/10i \$rip print the next 10 asm instructions

► live demo with `gdb_try_breaks.c`



If there are macros in your code, they can be expanded, provided that you compiled the code with the appropriate option:

```
-g3 [gdb3] [-dwarf-4]
```

(gdb) macro expand *macro*

shows the expansion of macro *macro*; *expression* can be any string of tokens

(gdb) info macro [-a|-all] *macro*

shows the current (or all) definition(s) of *macro*

(gdb) info macros *location*

shows all macro definitions effective at *location*



# Debugging: watch points



You can set *watchpoints* (aka “keep an eye on this and that”) instead of breakpoints, to stop the execution whenever a value of an expression / variable / memory region changes

(gdb) *watch variable*

keep an eye onto *variable*

(gdb) *watch expression*

stops when the value of expression changes (the scope of variables is respected)

(gdb) *watch -l expression*

Interpret *expression* as a memory location to be watched

(gdb) *watch -l expression*  
[mask *maskvalue*]

a mask for memory watching: specifies what bits of an address should be ignored (to match more addresses)

(gdb) *rwatch [-l] expr*  
[mask *mvalue*]

stops when the value of *expr* is read

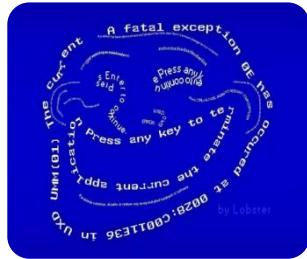
► live demo with `gdb_try_watch.c`



# Debugging outline



Introduction  
Debugging  
a code



Inspecting  
a code crash

Debugging  
a running  
process



It happens that you have code crashes in conditions not easily reproducible when you debug the code itself, for a number of reasons.

However, the O.S. can dump the entire “program status” on a file, called the *core file*:

```
luca@GGG:~/code/tricks% ./gdb_try_watch
no arguments were given, using default: 100

something wrong at point 2
[1] 8435 segmentation fault (core dumped) ./gdb_try_watch
luca@GGG:~/code/tricks% ls -l core
-rw----- 1 luca luca 413696 nov  8 15:45 core
luca@GGG:~/code/tricks% 
```

In order to allow it to dump the core, you have to check / set the core file size limit:

```
%> ulimit -c [size limit in KB]
```



Once you have a core, you can inspect it with GDB

```
%> gdb executable_name core_file_name
```

Or

```
%> gdb executable_name
(gdb) core core_file_name
```

The first thing to do, normally, is to unwind the stack frame to understand where the program crashed:

```
(gdb) bt full
```



# Debugging outline



Introduction



Debugging  
a code



Inspecting  
a code crash



Debugging  
a running  
process



In order to debug a running process, you can simply attach gdb to it:

%> gdb

(gdb) attach *process-id*

and start searching it to understand what is going on



# Debugging outline



GUI  
for GDB





1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS
5. DDD
6. NEMIVER, ECLIPSE, NETBEANS, CODEBLOCKS,  
many others..



You can start gdb with a text-user-interface:

```
%> gdb -tui
```

Or you can activate/deactivate it from gdb itself:

(gdb) ctrl-x a	
(gdb) ctrl-x o	
(gdb) ctrl-x 2	
(gdb) layout src	
asm	
split	
regs	

change focus
shows assembly windows
display src and commands
display assembly and commands
display src, asm and commands
display registers window



# GDB built-in tui



```
gdb_try_breaks.c
372  {
373
B+> 374      if ( argc > 1 )
375          // arg 0 is the name of the program itself
376          {
377              printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378              for ( int i = 1; i < argc; i++ )
379              {
380                  printf( "\targument %d is : %s\n", i, *(argv+i) );
381              }
382              printf( "\n" );
383          }
384
385      else
386
387          printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );
388
389
390      int arg1;
391
392      if ( argc > 1 )
393          arg1 = atoi( *(argv+1) );
394
395      else
396          arg1 = DEFAULT_ARG1;
397
398      int ret;
399
400      ret = function 1( arg1 );
```

```
native process 8943 In: main
(gdb) l
360    in /home/luca/code/tricks/gdb_try_breaks.c
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks

Breakpoint 1, main (argc=1, argv=0x7fffffffdaa8) at gdb_try_breaks.c:374
(gdb) 
```



# GDB built-in tui



```
gdb_try_breaks.c
B+> 374      if ( argc > 1 )
375          // arg 0 is the name of the program itself
376          {
377              printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378              for ( int i = 1; i < argc; i++ )
379              {
380                  printf( "\targument %d is : %s\n", i, *(argv+i) );
381              }
382              printf( "\n" );
383          }
384      else
385
386
387      printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );
```

```
0x555555554d38 <main+15>    cmpl    $0x1,-0x14(%rbp)
0x555555554d3c <main+19>    jle     0x555555554db7 <main+142>
0x555555554d3e <main+21>    cmpl    $0x2,-0x14(%rbp)
0x555555554d42 <main+25>    jle     0x555555554d4b <main+34>
0x555555554d44 <main+27>    mov     $0x73,%edx
0x555555554d49 <main+32>    jmp     0x555555554d50 <main+39>
0x555555554d4b <main+34>    mov     $0x20,%edx
0x555555554d50 <main+39>    mov     -0x14(%rbp),%eax
0x555555554d53 <main+42>    sub     $0x1,%eax
0x555555554d56 <main+45>    mov     %eax,%esi
0x555555554d58 <main+47>    lea     0x2bf(%rip),%rdi      # 0x5555555501e
0x555555554d5f <main+54>    mov     $0x0,%eax
0x555555554d64 <main+59>    callq   0x555555554680 <printf@plt>
0x555555554d69 <main+64>    movl   $0x1,-0xc(%rbp)
```

```
native process 9102 In: main
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks
Breakpoint 1, main (argc=1, argv=0xffffffffdaa8) at gdb_try_breaks.c:374
(gdb) layout split
```



GUI for  
GDB

# GDB dashboard



<https://github.com/cyrus-and/gdb-dashboard>

The screenshot shows the GDB dashboard interface with the following sections:

- Source:** Displays the C source code for a function named `fun` and its assembly translation.
- Assembly:** Shows the assembly code for the current instruction, with registers and memory references.
- Threads:** Lists two threads: thread 4355 and thread 0.
- Stack:** Displays the stack dump for both threads.
- Registers:** Shows the values of various CPU registers for both threads.
- Expressions:** Evaluates expressions like `data[i]` and `Memory`.
- Memory:** Displays memory dump for specific addresses.
- History:** Shows the history of commands entered.

```
>>> dashboard -output /dev/ttys001
>>> dashboard -layout
source
assembly
threads
stack
registers
expressions
memory
history
>>> p data[1]@2
$3 = {[0] = 0x7fff5fbffcf0 "hello", [1] = 0x7fff5fbffcf6 "GDB"}
>>> 
```



GUI for  
GDB

# GDBgui



<https://gdbgui.com/>

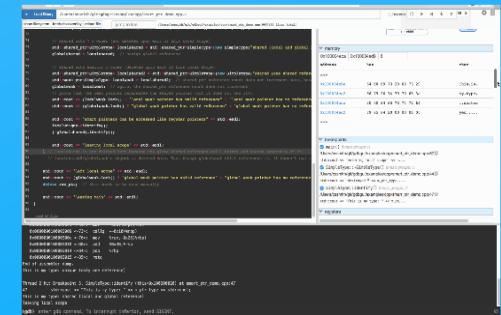
[gdbgui](#) Download Docs Examples Screenshots Videos About Chat GitHub

# gdbgui

Browser-based debugger for C, C++, go, rust, and more  
gdbgui turns this

```
>> gdb
GNU gdb (Ubuntu 7.11.1-0ubuntu1-16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This program is configured to run on i686-linux-gnu.
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
>>>
```

into this



[Download Now »](#)

[Watch Documentation](#)

that's all, have fun

"So long  
and thanks  
for all the fish"